Shape-CFD v5

AD-Rank & Shape-CFD: Physics-Inspired Reranking via Conservative Advection-Diffusion on Token-Level Point Clouds for Dense Retrieval

Yifan Chen

March 2026 (v5: Token-Level Point Clouds & PDE Orthogonal Fusion)

Abstract

Neural dense retrieval systems typically rank candidate documents by computing independent, pairwise cosine similarities between query and document embeddings, discarding the rich inter-document structural information that resides in the retrieved set. We present **Shape-CFD**, a physics-inspired reranking framework that recasts the post-retrieval reranking problem as a conservative convection-diffusion process on a sparse nearest-neighbor graph constructed over the candidate documents. Our formulation introduces two principal innovations: (1) an **advective transport term** that encodes query-directed momentum along graph edges, complementing the isotropic diffusion employed by prior graph-based rerankers; and (2) a **shape-aware extension** that upgrades each document from a single point embedding to a sentence-level point cloud, measuring inter-document affinity via the symmetric Chamfer distance. To mitigate the variance collapse of inner products in high-dimensional spaces (d=4096d = 4096), we apply Johnson-Lindenstrauss random projection to a compact m=128m = 128 subspace, which simultaneously achieves a 32×32\times computational speedup and introduces a stochastic regularization effect analogous to simulated annealing. The resulting iterative scheme provably converges to a unique stationary distribution under a Courant-Friedrichs-Lewy (CFL) stability condition, and its conservative upwind discretization guarantees global mass conservation. On three BEIR benchmarks, Shape-CFD achieves competitive NDCG@10 scores—including a +**29.9%** improvement on NFCorpus (NDCG@10 = **0.2852**, all-time best via the calibrated

Rust full-pipeline engine)—while maintaining a retrieval latency of **38 ms** on commodity CPU hardware via a native Rust full-pipeline engine, approximately 13×13\times faster than cross-encoder rerankers. **In this v3 supplement**, we report a systematic ablation of the Allen-Cahn reaction-diffusion extension (proposed as future work in v1) across NFCorpus and SciFact, demonstrate that nonlinear reaction dynamics are numerically dormant in the narrow score-range regime of short queries (ΔC0≈0.1\Delta C_0 \approx 0.1), and discover that a mixed initial field C0=α·MaxSim+(1−α)·MeanSimC_0 = \alpha \cdot \text{MaxSim} + (1-\alpha) \cdot \text{MeanSim} with α=0.7\alpha = 0.7 yields a +**3.2%** improvement over Shape-CFD v4 on NFCorpus—but degrades on SciFact. We formalize this dataset-dependent behavior as the **Pareto boundary of initial fields**, governed by the trade-off between extreme-value sampling noise and signal dilution, and propose a scale-adaptive symmetry-breaking formula that reduces to the original Chamfer baseline as a special case. Additionally, we introduce **PQ-Chamfer** (Subspace Chamfer Distance), a structured decomposition that partitions the 4096-dimensional embedding space into 64 contiguous 64-dimensional subspaces and computes independent cosine distances within each, yielding a +**8.8%** improvement over the full-dimensional baseline—compared to the −7.6% degradation of random JL projection—demonstrating that preserving dimensional structure is critical for high-dimensional distance computation. **In this v4 supplement**, we discover that treating each PQ subspace as an independent "virtual token" and reversing the order of subspace aggregation and sentence-level Chamfer computation (VT-Aligned) yields an additional +**2.5%** improvement at zero cost—no model re-inference, no new data. Combined with the V7.1 adjoint state prefetch method, the stacked **VT-V7** configuration achieves +29.6% over the cosine baseline on NFCorpus (NDCG@10 = 0.2844) in JavaScript. A calibrated native Rust implementation with 6 bug fixes further improves to +**29.9%** (NDCG@10 = **0.2852**), with end-to-end latency of **38 ms** per query. **In this v5 supplement**, we introduce **token-level point clouds**: by extracting per-token last hidden states from the embedding model (rather than pooled sentence embeddings), each document becomes a dense point cloud of $$356 tokens and each query unfolds into $$6 token vectors, resolving the "single-point query degeneracy" that limited prior sentence-level formulations. The **Token PQ-Chamfer** distance—a subspace-decomposed symmetric Chamfer computed over token-level point clouds—achieves NDCG@10 = **0.3214** (+46.4% vs cosine) via full scan, and a **two-stage centroid acceleration** (centroid coarse filter top-100 followed by full token Chamfer re-ranking) reaches **0.3220** (+46.7%) in only **23 ms** with lossless accuracy. Finally, we discover that PDE-based reranking serves as an **orthogonal signal source** rather than a same-perspective smoother at the token level: score-level fusion s=λ·sTC+(1−λ)·sPDEs = \lambda \cdot s_{TC} + (1-\lambda) \cdot s_{PDE} with λ=0.7\lambda = 0.7 achieves the overall best NDCG@10 = **0.3232**

(+47.2% vs cosine) at **26 ms** latency on CPU, with zero cross-encoder inference and zero training.

---

1 Introduction

Retrieval-Augmented Generation (RAG) systems have emerged as the dominant architecture for grounding large language models in external knowledge [10, 11]. In a typical RAG pipeline, a bi-encoder first retrieves a set of NN candidate passages from a large corpus using approximate nearest-neighbor (ANN) search, and then either these candidates are fed directly to the generator or an intermediate reranker adjusts the ordering before generation. The quality of the final answer critically depends on the precision and recall of this retrieval stage.

The prevailing approach to first-stage ranking relies on cosine similarity between the query vector qq and each document vector viv_i, computed independently: score(di)=q·vi∥q∥ ∥vi∥ \text{score}(d_i) = \frac{q \cdot v_i}{\|q\| \|v_i\|} While simple and efficient, this pointwise scoring paradigm fundamentally ignores the **inter-document semantic structure** within the retrieved set. When the query is ambiguous, polysemous, or spans multiple topical domains, isolated scoring fails to exploit the mutual information carried by document neighborhoods—a limitation well-studied under the cluster hypothesis in information retrieval [4].

Cross-encoder rerankers [11] address this quality gap by jointly encoding the query-document pair through a deep transformer, achieving state-of-the-art ranking accuracy. However, their computational cost scales as O(N·L2·H)O(N \cdot L^2 \cdot H), where LL is the sequence length and HH the hidden dimension, resulting in latencies of approximately 500 ms for N=30N = 30 candidates—prohibitive for many real-time applications.

Recent work has explored graph-based diffusion as a middle ground. Dampanaboina et al. [2] construct a document similarity graph and apply isotropic heat diffusion to re-score candidates, demonstrating that inter-document signal propagation can improve retrieval quality without the overhead of cross-encoders. However, their formulation is purely diffusive: every node diffuses its score to neighbors symmetrically, without directional preference. In physical terms, this corresponds to a system governed solely by Fick’s second law, which guarantees convergence to a flat equilibrium—the retrieval analogue of over-smoothing in graph neural networks [13, 14].

Our key insight is that retrieval reranking is more faithfully modeled by the full convection-diffusion equation rather than pure diffusion. In fluid dynamics, the

convection-diffusion equation $\frac{\partial C}{\partial t} = D\nabla^2 C - \nabla \cdot (\vec{u}C)$ governs the transport of a scalar concentration field C under the combined influence of diffusive spreading (coefficient D) and directional advective transport (velocity field $\vec{u}$). By drawing this analogy, we introduce an **advection term** that encodes the query's semantic direction: on each graph edge, the advection coefficient measures the alignment between the inter-document displacement vector and the query embedding, providing a directional bias that pure diffusion lacks.

A second limitation of existing approaches is that they represent each document as a single point in the embedding space. In practice, a document—particularly a statutory provision or a technical passage—carries heterogeneous semantic content across its constituent sentences. A single centroid embedding necessarily conflates these distinct semantic facets. To address this, we propose the **Shape-CFD extension**, which represents each document as a sentence-level point cloud $\mathcal{P}_i = \{s_{i,1}, \ldots, s_{i,K_i}\} \subset \mathbb{R}^d$ and replaces cosine-based graph weights with the symmetric Chamfer distance. This formulation captures local semantic matches between document sub-components—even when the overall document-level similarity is moderate—and naturally degrades to the point representation for single-sentence documents.

However, computing Chamfer distances in the native $d = 4096$ embedding space introduces a dual challenge: (i) the computational cost of $O(K_A \cdot K_B \cdot d)$ per document pair, and (ii) the curse of dimensionality, which causes inner-product-based advection coefficients to collapse toward zero (with standard deviation $\sigma \approx 1/\sqrt{d} \approx 0.016$). We resolve both issues via Johnson-Lindenstrauss (JL) random projection to $m = 128$ dimensions, achieving a $32\times$ speedup. Crucially, we provide a corrected theoretical account of the dimensional reduction effect: contrary to naive interpretations, the expanded variance of inner products in the reduced space does not "amplify" true physical signals but rather injects $O(1/m)$ stochastic noise that acts as a form of flow-field annealing, breaking the near-orthogonal deadlock of high-dimensional geometry.

Contributions

We summarize our contributions as follows:

1. **AD-Rank**: A graph-based reranking algorithm that solves the conservative convection-diffusion equation on a sparse KNN document graph, introducing directional query-aware advection to complement isotropic

diffusion. We prove global mass conservation and convergence under CFL stability conditions.

2. **Shape-CFD**: An extension that replaces point embeddings with sentence-level point clouds, using Chamfer distance for graph construction and centroid-based advection directions. This captures fine-grained sub-document semantic matching.

3. **Theoretical analysis**: We characterize the operating regime through the structural Péclet number and provide a physically grounded explanation for the empirically observed optimality of weak advection (Pe≈0.065\text{Pe} \approx 0.065).

4. **VT-Aligned virtual token decomposition** (v4): By reversing the order of subspace aggregation and sentence-level Chamfer computation, we transform 64 PQ subspaces into independent "virtual tokens" that increase effective query density 64-fold at zero cost (+2.5% on NFCorpus).

5. **Adjoint state prefetch** (v4): A speculative retrieval mechanism that uses PDE concentration gradients to identify and expand the candidate pool beyond the initial retrieval horizon (+1.1% independently, +4.2% stacked with VT-Aligned).

6. **Native Rust full-pipeline engine** (v4): A complete retrieval-to-reranking pipeline that abandons HNSW and cosine recall in favor of fullscan + VT-Aligned distance + PDE reranking, achieving 38 ms/query with 45 unit tests.

7. **Comprehensive evaluation**: We evaluate on three BEIR benchmarks (SciFact, NFCorpus, FiQA) and a Chinese legal statute retrieval dataset (23,701 documents), demonstrating NDCG@10 = 0.2852 (+29.9% vs cosine) on NFCorpus with 38 ms latency on CPU.

8. **Token-level point clouds** (v5): By extracting per-token last hidden states (rather than pooled embeddings), each document becomes a dense point cloud of $$356 tokens and each query unfolds into $$6 token vectors, enabling fine-grained token-to-token matching that approximates hard cross-attention with zero parameters and zero training. Token PQ-Chamfer achieves NDCG@10 = 0.3214 (+46.4% vs cosine) on NFCorpus.

9. **Two-stage centroid acceleration** (v5): A lossless two-stage retrieval pipeline—token centroid coarse filtering (top-100) followed by full token PQ-Chamfer re-ranking—achieves a 13x speedup (305 ms to 23 ms) with no accuracy degradation (0.3220 vs 0.3214 full scan).

10. **PDE orthogonal fusion** (v5): We discover that PDE-based reranking provides an orthogonal signal (global semantic propagation) to token Chamfer (local nearest-neighbor matching). Score-level fusion with λ=0.7\lambda = 0.7 achieves the overall best NDCG@10 = **0.3232** (+47.2% vs cosine) at 26 ms latency.

The remainder of this paper is organized as follows. Section 2 surveys related work. Section 3 presents the mathematical framework of AD-Rank, its Shape-CFD extension, the VT-Aligned virtual token decomposition, the adjoint state prefetch method, and the token-level point cloud formulation with PDE fusion. Section 4 reports experimental results including the Rust full-pipeline engine and V11 token-level ablations. Section 5 provides physical analysis of the algorithm’s behavior, and Section 6 concludes with directions for future work.

---

2 Related Work

Our work draws upon and bridges three distinct research threads: graph-based reranking in information retrieval, distributional and shape-aware representations for text, and partial differential equations on graph structures.

2.1 Graph-Based Reranking

The idea of leveraging inter-document relationships for reranking has a long history. The PageRank algorithm [4] pioneered the use of iterative graph propagation for web page ranking, treating hyperlink structure as a random walk on a directed graph. More recently, several works have adapted graph-based methods specifically for dense retrieval reranking.

Diffusion-Aided RAG [2] constructs a semantic similarity graph over retrieved documents and applies isotropic heat diffusion to redistribute relevance scores. By allowing high-scoring documents to propagate their scores to semantically similar neighbors, this approach effectively discovers latent relevant passages that might have received low initial cosine scores. However, the diffusion process is purely isotropic—it lacks any directional preference tied to the query—and therefore cannot distinguish between semantically similar documents that differ in their query-relevance orientation.

G-RAG [18] employs graph neural networks to learn reranking functions over document graphs, achieving strong results through supervised training on labeled query-document relevance pairs. While powerful when sufficient training data is available, this supervised approach requires domain-specific annotation and does not transfer freely across domains without fine-tuning.

Graph-Augmented Dense Statute Retrieval (G-DSR) [19] exploits the hierarchical structure inherent in legal document collections, using graph augmentation to improve statute retrieval. Our approach differs fundamentally in that we construct graphs dynamically from the retrieved set at query time, making no assumptions about pre-existing document structure.

Spreading Activation for KG-RAG [20] adapts the classical spreading activation paradigm from cognitive science to knowledge graph-augmented retrieval. Activation propagation can be viewed as a special case of diffusion; our framework generalizes this by introducing the advection term for directional transport.

2.2 Shape and Distributional Representations

The standard practice in dense retrieval compresses each document into a single vector (a "point" in embedding space), which necessarily conflates the diverse semantic facets of multi-sentence passages. Several prior works have explored richer representations.

Gaussian Embeddings [1] represent words as Gaussian distributions rather than point vectors, capturing uncertainty and enabling asymmetric distance measures via KL divergence. The GELD framework [21] extends this to document-level Gaussian embeddings for retrieval.

ColBERT [22] introduces a late interaction mechanism in which both query and document are represented as sets of token-level vectors, and relevance is computed via a MaxSim operator that finds the maximum similarity between each query token and all document tokens. This multi-vector representation captures fine-grained alignment but computes relevance in a purely pointwise manner—each document is scored independently of its retrieved neighbors.

Our Shape-CFD extension shares ColBERT's insight that documents should be represented as collections of sub-vectors. However, we go further by using these point clouds not for isolated pointwise scoring but as the basis for graph-structured collective reranking: the Chamfer distance between point clouds determines graph edge weights, and the full convection-diffusion dynamics then propagates information across the document graph. This listwise, topology-aware approach is fundamentally distinct from ColBERT's pointwise MaxSim.

2.3 Partial Differential Equations on Graphs

The application of continuous PDEs to discrete graph structures has gained significant traction in the graph neural network (GNN) community.

GRAND [13] established that many GNN architectures can be interpreted as discretizations of the heat diffusion equation on graphs, providing a continuous-time perspective on message passing. This insight has spawned a family of diffusion-based GNN architectures.

Reaction-Diffusion GNNs [23] extend the pure diffusion framework by incorporating nonlinear reaction terms, enabling the formation of Turing patterns—spatially heterogeneous steady states that counteract the over-smoothing tendency of pure diffusion. The Allen-Cahn reaction term, in particular, drives a phase-separation dynamic that polarizes node features into distinct clusters. Our analysis of the Péclet number regime (Section 5) reveals a related phenomenon: the dominant diffusion term provides global smoothing while the weak advection term introduces a subtle but critical directional perturbation.

Advection-Augmented CNNs [3] directly incorporate advection-diffusion dynamics into convolutional neural network layers for scientific prediction tasks, demonstrating that the advective transport component meaningfully enhances model performance. Our work is, to our knowledge, the first to apply the full convection-diffusion equation to information retrieval reranking.

Spectral methods on graphs, including Chebyshev spectral filtering [24] and diffusion kernels via the matrix exponential $K=\exp(-\beta L)$ [12], provide alternative computational pathways. For our problem scale ($N=30$ candidate documents), explicit time-stepping proves more practical than eigendecomposition, though spectral acceleration remains a promising direction for larger candidate sets.

---

3 Method

We now present the AD-Rank algorithm and its Shape-CFD extension. We begin by formalizing the reranking problem, then describe the graph construction, the conservative convection-diffusion discretization, convergence guarantees, and the shape-aware generalization.

3.1 Problem Formulation

Let $q \in \mathbb{R}^d$ denote the query embedding produced by a pretrained language model (in our implementation, Qwen3-Embedding-8B with $d=4096$). Given a corpus, the first-stage retriever returns a candidate set $\mathcal{D} = \{d_1, d_2, \ldots, d_N\}$ of N documents (we use $N=30$ via HNSW approximate nearest-neighbor search), where each

document did_i is associated with an embedding vector vi∈ℝdv_i \in \mathbb{R}^d.

The reranking task is to learn a scoring function f:𝒟×q→ℝNf: \mathcal{D} \times q \to \mathbb{R}^N that produces a ranking superior to the initial cosine-similarity ordering. Unlike pointwise and pairwise rerankers, our approach is listwise: the score of each document is determined collectively by the entire retrieved set through a dynamic process on a document graph.

3.2 Graph Construction

We construct a sparse, weighted, undirected graph 𝒢=(𝒱,ℰ,W)\mathcal{G} = (\mathcal{V}, \mathcal{E}, W) over the NN candidate documents.

Step 1: Pairwise similarity. We compute the full N×NN \times N cosine similarity matrix: Sij=cos(vi,vj)=vi·vj∥vi∥∥vj∥,∀i,j∈{1,⋯,N}S_{ij} = \cos(v_i, v_j) = \frac{v_i \cdot v_j}{\|v_i\| \|v_j\|}, \quad \forall i,j \in \{1, \ldots, N\}

Step 2: KNN sparsification. For each node ii, we retain edges only to its kk most similar neighbors: 𝒩(i)=argTopKj≠i(Sij,k),k=3\mathcal{N}(i) = \text{argTopK}_{j \neq i}(S_{ij}, k), \quad k = 3 This sparsification reduces the graph density from O(N2)O(N^2) to O(Nk)O(Nk), which is critical for controlling both computational cost and diffusion dynamics.

Step 3: Edge weights. We assign exponentially decaying weights: Wij=exp(−β·(1−Sij)),β=2.0W_{ij} = \exp\left(-\beta \cdot (1 - S_{ij})\right), \quad \beta = 2.0 The parameter β\beta controls the decay rate: larger values concentrate diffusion along high-similarity edges.

Step 4: Symmetrization. If j∈𝒩(i)j \in \mathcal{N}(i) but i∉𝒩(j)i \notin \mathcal{N}(j), we add the reverse edge with Wji=WijW_{ji} = W_{ij}. This ensures the graph is undirected, which is essential for the symmetric positive-definite property of the graph Laplacian and, consequently, for the conservation law of the diffusion operator.

3.3 Conservative Convection-Diffusion on Graphs

We model the evolution of document relevance scores as a convection-diffusion process on the graph 𝒢\mathcal{G}.

Initial concentration field. Each document is assigned an initial "concentration" equal to its cosine similarity with the query:

C0(i)=cos(q,vi),i=1,⋯,NC_0(i) = \cos(q, v_i), \quad i = 1, \ldots, N This initial condition seeds the dynamic process with the first-stage retrieval signal.

Advection direction. For each edge (i,j)(i, j), we define the advection coefficient by projecting the inter-document displacement onto the query direction: uij=⟨vj−vi∥vj−vi∥2+ε,q∥q∥⟩⋅ustrength,ε=10−8u_{ij} = \left\langle \frac{v_j - v_i}{\|v_j - v_i\|_2 + \varepsilon},\ \frac{q}{\|q\|} \right\rangle \cdot u_{\text{strength}}, \quad \varepsilon = 10^{-8}

The scalar uiju_{ij} measures how well the edge direction i→ji \to j aligns with the query’s semantic orientation. A positive value (uij>0u_{ij} > 0) indicates that moving from ii to jj proceeds “downwind” (toward the query); a negative value indicates the reverse. The small constant ε\varepsilon prevents division by zero when two documents have near-identical embeddings.

Anti-symmetry. The advection coefficients satisfy a strict anti-symmetry: uji=⟨vi−vj∥vi−vj∥,q⟩⋅ustrength=−uiju_{ji} = \left\langle \frac{v_i - v_j}{\|v_i - v_j\|},\ \hat{q} \right\rangle \cdot u_{\text{strength}} = -u_{ij} since vi−vj=−(vj−vi)v_i - v_j = -(v_j - v_i) and ∥vi−vj∥=∥vj−vi∥\|v_i - v_j\| = \|v_j - v_i\|. This property is crucial for ensuring global conservation of the advective flux (∑iFiadv=0\sum_i F_i^{\text{adv}} = 0).

Governing equation. The continuous convection-diffusion equation ∂C∂t=D∇2C−∇⋅(u→C)\frac{\partial C}{\partial t} = D \nabla^2 C - \nabla \cdot (\vec{u} C) is discretized on the graph using a **conservative upwind scheme**:

Ct+1(i)=Ct(i)+Δt⋅[D∑j∈𝒩(i)Wij(Ct(j)−Ct(i))⏟diffusion flux+∑j∈𝒩(i)Wij[max(uji,0)Ct(j)−max(uij,0)Ct(i)]⏟conservative upwind advection flux]C_{t+1}(i) = C_t(i) + \Delta t \cdot \Bigg[\underbrace{D \sum_{j \in \mathcal{N}(i)} W_{ij}\, \big(C_t(j) - C_t(i)\big)}_{\text{diffusion flux}} + \underbrace{\sum_{j \in \mathcal{N}(i)} W_{ij} \Big[\max(u_{ji}, 0)\, C_t(j) - \max(u_{ij}, 0)\, C_t(i)\Big]}_{\text{conservative upwind advection flux}} \Bigg]

The upwind discretization is essential for numerical stability. A naive central-difference scheme (uijWij(Cj−Ci)u_{ij} W_{ij} (C_j - C_i)) produces non-physical oscillations when the local Péclet number exceeds 2. The upwind scheme resolves this by always sampling concentration from the upstream node, preserving the causal direction of information transport.

Global conservation. We prove that the total concentration is conserved at each time step:

∑i=1NFiadv=∑i=1N∑j∈𝒩(i)Wij[max(uji,0)Cj−max(uij,0)Ci]=0\sum_{i=1}^{N} F_i^{\text{adv}} = \sum_{i=1}^{N} \sum_{j \in \mathcal{N}(i)} W_{ij} \left[\max(u_{ji}, 0)\, C_j - \max(u_{ij}, 0)\, C_i\right] = 0

Proof. Since the graph is undirected (Wij=WjiW_{ij} = W_{ji}) and uji=−uiju_{ji} = -u_{ij}, each ordered pair (i,j)(i, j) contributes a flux of Wijmax(uji,0)CjW_{ij} \max(u_{ji}, 0) C_j to node ii and Wjimax(uij,0)CiW_{ji} \max(u_{ij}, 0) C_i outward from node jj (equivalently, −Wijmax(uij,0)Ci-W_{ij} \max(u_{ij}, 0) C_i at node ii). Summing over all edges, each pair’s contributions cancel exactly. ■\blacksquare

This conservation property guarantees that the reranking process redistributes relevance scores without creating or destroying total “score mass”—a desirable property that prevents score inflation or deflation.

CFL stability condition. To ensure absolute stability of the explicit time integration (non-negative concentrations and absence of numerical oscillations), we require that the coefficient of Ct(i)C_t(i) in the update equation remains non-negative: 1−Δt∑j∈𝒩(i)Wij[D+max(uij,0)]≥01 - \Delta t \sum_{j \in \mathcal{N}(i)} W_{ij} \left[D + \max(u_{ij}, 0)\right] \geq 0

This yields the strict CFL time-step bound: ΔtCFL=1maxi∑j∈𝒩(i)Wij(D+max(uij,0))\Delta t_{\text{CFL}} = \frac{1}{\max_{i} \sum_{j \in \mathcal{N}(i)} W_{ij} \left(D + \max(u_{ij}, 0)\right)}

In practice, we use Δt=Csafe·ΔtCFL\Delta t = C_{\text{safe}} \cdot \Delta t_{\text{CFL}} with safety factor Csafe=0.9C_{\text{safe}} = 0.9. The iteration terminates when maxi|Ct+1(i)−Ct(i)|<ϵ\max_i |C_{t+1}(i) - C_t(i)| < \epsilon with ϵ=10−3\epsilon = 10^{-3}, which typically requires 30–50 steps.

3.4 Convergence Analysis

Theorem 1 (Global Convergence). Given a connected graph 𝒢\mathcal{G} and a time step Δt\Delta t satisfying the CFL condition of Section 3.3, the iterative sequence {Ct}\{C_t\} generated by the conservative upwind scheme converges to a unique stationary distribution C*C^*, and the total concentration ∑iCt(i) \sum_i C_t(i) is conserved for all tt.

Proof. We write the update in matrix form Ct+1=**M**CtC_{t+1} = \mathbf{M}\, C_t, where the transition matrix **M**\mathbf{M} has entries: - Diagonal: **M**ii=1−Δt∑k∈𝒩(i)Wik[D+max(uik,0)]\mathbf{M}_{ii} = 1 - \Delta t \sum_{k \in \mathcal{N}(i)} W_{ik} \left[D + \max(u_{ik}, 0)\right] - Off-diagonal (j∈𝒩(i)j \in

$\mathcal{N}(i)$): $\mathbf{M}_{ij} = \Delta t \cdot W_{ij} \left[D + \max(u_{ji}, 0)\right]$

We establish the four conditions for the Perron-Frobenius theorem:

(i) Non-negativity. The CFL condition ensures $\mathbf{M}_{ii} \geq 0$; all off-diagonal entries are non-negative since $D, W_{ij} \geq 0$.

(ii) Column-stochasticity. By the symmetry of edge weights ($W_{ij} = W_{ji}$) and the structure of the upwind scheme, the column sums satisfy $\sum_i \mathbf{M}_{ij} = 1$ for all j. Thus $\mathbf{M}$ is column-stochastic, and the total concentration is conserved.

(iii) Irreducibility. When the graph $\mathcal{G}$ is connected, $\mathbf{M}$ is irreducible.

(iv) Aperiodicity. Since $\mathbf{M}_{ii} > 0$ for all i (the CFL bound is strict), every node has a self-loop, guaranteeing aperiodicity.

By the Perron-Frobenius theorem, the dominant eigenvalue $\lambda_1 = 1$ is simple, and all other eigenvalues satisfy $|\lambda_k| < 1$. Therefore, $C_t = \mathbf{M}^t C_0$ converges exponentially to the unique stationary distribution C^*. $\blacksquare$

Remark on graph connectivity. In practice, a KNN graph with $N = 30$ and $k = 3$ may decompose into multiple connected components, particularly for cross-domain queries that fragment the candidate set into semantic islands. When the graph has $K > 1$ connected components, the eigenvalue $\lambda_1 = 1$ has algebraic multiplicity K, and the stationary distribution is no longer unique—it depends on the initial allocation of concentration mass across components. We note that this is, in fact, a desirable property: it naturally prevents score leakage between unrelated semantic clusters. In our implementation, we do not enforce artificial connectivity (e.g., via teleportation as in PageRank), allowing each semantic sub-manifold to evolve independently.

Early stopping and transient dynamics. The fully converged stationary C^* depends only on the graph topology and advection field, and is independent of the initial condition C_0. In practice, we terminate at $\epsilon = 10^{-3}$ (typically 30–50 iterations), capturing the transient state where the initial retrieval signal (C_0) has been partially—but not completely—smoothed by the diffusion dynamics. This transient truncation is not merely an engineering convenience but a deliberate design choice: it preserves the first-stage cosine relevance signal while enriching rankings with inter-document contextual information.

3.5 Adaptive Parameters

Fiedler-value adaptive diffusion. The algebraic connectivity of the graph (the second-smallest eigenvalue λ2\lambda_2 of the graph Laplacian $L = D_{} - W, known as the Fiedler value) provides a natural measure of graph density. We adapt the diffusion coefficient accordingly: Dadaptive=Dbase1+λ2D_{\text{adaptive}} = \frac{D_{\text{base}}}{1 + \lambda_2} When λ2\lambda_2 is large (tightly clustered candidates), DD is reduced to prevent over-smoothing; when λ2\lambda_2 is small (sparse, dispersed candidates), DD is increased to promote signal propagation.

Cross-domain adaptive advection. The variance of inter-document cosine similarities σsim2\sigma^2_{\text{sim}} serves as a proxy for semantic heterogeneity. We adjust the advection strength via linear interpolation: uadaptive=umin+(umax−umin)⋅min(σsim20.04,1)u_{\text{adaptive}} = u_{\min} + (u_{\max} - u_{\min}) \cdot \min\left(\frac{\sigma^2_{\text{sim}}}{0.04},\ 1\right) Higher variance (indicating cross-domain queries) triggers stronger advection to exploit the more pronounced directional gradients.

3.6 Shape-CFD Extension

The core innovation of Shape-CFD is the elevation of document representations from single points to multi-point shapes (point clouds).

Point cloud representation. Each document did_i is decomposed into its constituent sentences, and each sentence is independently embedded to form a point cloud: di→𝒫i={si,1,si,2,⋯,si,Ki}⊂ℝdd_i \to \mathcal{P}_i = \{s_{i,1}, s_{i,2}, \ldots, s_{i,K_i}\} \subset \mathbb{R}^d where KiK_i is the number of sentences in document did_i. Single-sentence documents naturally reduce to singleton point clouds, ensuring backward compatibility with the point-based AD-Rank.

Chamfer distance as Unbalanced Optimal Transport. We measure inter-document affinity using the symmetric Chamfer distance: dChamfer(𝒫A,𝒫B)=1|𝒫A|∑a∈𝒫Aminb∈𝒫Bdcos(a,b)+1|𝒫B|∑b∈𝒫Bmina∈𝒫Adcos(a,b)d_{\text{Chamfer}}(\mathcal{P}_A, \mathcal{P}_B) = \frac{1}{|\mathcal{P}_A|} \sum_{a \in \mathcal{P}_A} \min_{b \in \mathcal{P}_B} d_{\cos}(a, b) + \frac{1}{|\mathcal{P}_B|} \sum_{b \in \mathcal{P}_B} \min_{a \in \mathcal{P}_A} d_{\cos}(a, b) where dcos(a,b)=1−cos(a,b)d_{\cos}(a, b) = 1 - \cos(a, b) is the cosine distance. Unlike the centroid cosine similarity, Chamfer captures local semantic matches: even when two documents are globally dissimilar, high similarity between specific sentence pairs yields a small Chamfer distance.

Crucially, Chamfer distance is not an approximation to the Wasserstein-1 (Earth Mover’s) distance, but rather a fundamentally different object. The standard W_1 distance enforces strict mass conservation via marginal constraints: every unit of semantic mass in the source must be transported to a unique location in the target. For retrieval, this creates a catastrophic noise penalty: when a query contains 1 sentence and a document contains 50 sentences (only 1 relevant), W_1 forces the 49 irrelevant sentences to be matched somewhere, generating spurious transport cost that overwhelms the genuine relevance signal.

Chamfer distance operates as an Unbalanced Optimal Transport (UOT) without marginal constraints. It allows many-to-one mappings and abandons unmatched points entirely, functioning as a parameter-free hard semantic attention mechanism that isolates the most relevant sentence pairs while discarding noise. This property makes Chamfer distance particularly well-suited for long, heterogeneous documents—precisely where Shape-CFD provides the greatest advantage (NFCorpus: +24.6% vs. cosine).

Shape-aware graph weights and initial conditions. We replace the cosine-based graph construction with Chamfer-based equivalents:
$W_{ij}^{\text{shape}} = \exp\left(-\beta \cdot d_{\text{Chamfer}}(\mathcal{P}_i, \mathcal{P}_j)\right)$
$C_0(i) = \exp\left(-2 \cdot d_{\text{Chamfer}}(\mathcal{P}_q, \mathcal{P}_i)\right)$

The exponential mapping ensures non-negative initial concentrations ($C_0 \geq 0$), resolving the physical inconsistency of negative concentrations that arises when using raw cosine similarities (which can be negative) as initial conditions in the diffusion process.

Centroid-based advection. For the advection direction, we compute the normalized centroid of each point cloud:
$\bar{v}_i = \frac{\sum_{k=1}^{K_i} s_{i,k}}{\left\| \sum_{k=1}^{K_i} s_{i,k}\right\|_2}$ and define the shape-aware advection coefficient as: $u_{ij}^{\text{shape}} = \left\langle \frac{\bar{v}_j - \bar{v}_i}{\|\bar{v}_j - \bar{v}_i\|}, \frac{\bar{v}_q}{\|\bar{v}_q\|} \right\rangle \cdot u_{\text{strength}}$

The explicit L_2-normalization of the centroid sum is critical: without it, documents with many semantically diverse sentences would have centroids that collapse toward the origin ($\|\bar{v}_i\| \to 0$), introducing a severe document-length bias into the advection field.

Metric space inconsistency (known limitation). We acknowledge a fundamental tension in this formulation: the graph topology is constructed using Chamfer distance (a set-theoretic metric on point clouds), while the advection coefficients are computed from centroid differences (a Euclidean metric on single points). This places the diffusion and advection operators in two different metric spaces—a gap that is theoretically unsatisfying. We adopt this hybrid approach as a pragmatic engineering compromise in the current version: Chamfer distance captures fine-grained inter-document structure, while centroid-based advection remains computationally tractable. In Section 6, we propose scalar potential convection as a principled resolution that eliminates this inconsistency entirely.

Noise dilution in symmetric Chamfer (known limitation). The symmetric Chamfer distance averages over both forward and reverse nearest-neighbor distances. In retrieval scenarios where a long document (Ki=50K_i = 50 sentences) contains only a single highly relevant sentence, the remaining 49 unmatched sentences contribute large cosine distances to the reverse term, diluting the signal from the one critical match. This noise dilution effect systemically penalizes long, heterogeneous documents—precisely the documents where shape-aware representations should provide the greatest benefit. Asymmetric alternatives, such as the directed Hausdorff distance or ColBERT-style MaxSim, would address this issue at the cost of losing the metric symmetry required by our undirected graph framework. We defer the formal resolution to future work (Section 6).

3.7 Acceleration via Johnson-Lindenstrauss Projection

Computing Chamfer distances in the native d=4096d = 4096 space has complexity O(KA·KB·d)O(K_A \cdot K_B \cdot d) per document pair. For N=30N = 30 candidates with an average of K=5K = 5 sentences each, this amounts to (302)×25×4096≈4.5×107\binom{30}{2} \times 25 \times 4096 \approx 4.5 \times 10^7 multiply-adds—feasible but expensive for real-time applications.

We apply a random projection P∈ℝm×dP \in \mathbb{R}^{m \times d} with Pij∼𝒩(0,1/m)P_{ij} \sim \mathcal{N}(0, 1/m) and m=128m = 128:
s̃=P·s∥P·s∥2\tilde{s} = \frac{P \cdot s}{\|P \cdot s\|_2}

The post-projection L2L_2-normalization ensures that the equivalence ∥x−y∥22=2(1−cos(x,y))\|x - y\|_2^2 = 2(1 - \cos(x, y)) holds for unit vectors, allowing the JL distance-preservation guarantee to transfer to the cosine metric.

Johnson-Lindenstrauss Lemma [7]. For any ϵ>0\epsilon > 0 and point set of size nn, when m⩾O(ϵ−2logn)m \geq O(\epsilon^{-2} \log n):

(1−ϵ)∥sa−sb∥2≤∥s̃a−s̃b∥2≤(1+ϵ)∥sa−sb∥2(1 - \epsilon)\|s_a - s_b\|^2 \leq \|\tilde{s}_a - \tilde{s}_b\|^2 \leq (1 + \epsilon)\|s_a - s_b\|^2

Dimensional reduction effects and their limitations. In the reduced m=128m = 128 space, the standard deviation of random inner products increases from σ=1/4096≈0.016\sigma = 1/\sqrt{4096} \approx 0.016 to σ=1/128≈0.088\sigma = 1/\sqrt{128} \approx 0.088—a factor of ∼5.6×\sim 5.6\times. The JL lemma guarantees that the expected value of pairwise distances is preserved (𝔼[∥s̃a−s̃b∥2]=∥sa−sb∥2\mathbb{E}[\|\tilde{s}_a - \tilde{s}_b\|^2] = \|s_a - s_b\|^2), so the graph topology constructed from Chamfer distances remains faithful in expectation. However, the expanded variance of inner products in the reduced space is projection noise, not signal amplification. This noise introduces O(1/m)O(1/m) stochastic perturbations into the advection coefficients.

We make no claim that this noise is physically beneficial. The advection coefficients in the native d=4096d = 4096 space are inherently weak due to the concentration of measure on the unit sphere (σ≈0.016\sigma \approx 0.016), and the JL projection does not resolve this fundamental limitation—it merely trades the curse of dimensionality for projection variance. The empirical effectiveness of the algorithm in the projected space should be attributed to the diffusion-dominated regime (Pe ≈0.065\approx 0.065, where advection plays a secondary role) rather than to any purported "signal amplification" of the projection. As discussed in Section 6, scalar potential convection—which operates on scalar score differences rather than vector inner products—offers a principled alternative that is immune to the curse of dimensionality.

Computational speedup. The primary justification for JL projection is computational: the dimensionality reduction yields a d/m=32×d/m = 32\times speedup in all distance computations. Combined with offline precomputation and SQLite caching of projected sentence vectors (a one-time cost), the end-to-end Shape-CFD latency is reduced from 1,837 ms (with live API embedding) to **22 ms**.

3.7.1 VT-Aligned: Virtual Token Decomposition (v4)

The PQ-Chamfer decomposition (Section 4.3.5) partitions the 4096-dimensional embedding space into S=64S = 64 contiguous subspaces of dimension ds=64d_s = 64 and computes independent cosine distances within each. However, the order of operations—subspace aggregation versus sentence-level Chamfer—critically affects retrieval quality. We discover that reversing the standard order unlocks an additional +2.5% improvement at zero computational cost.

Standard PQ-Chamfer (subspace-first). The original PQ-Chamfer first aggregates across SS subspaces to produce a single scalar distance between two vectors, then applies the Chamfer operator across sentences: dPQ-Chamfer(q,D)=Chamfer({q},D;dPQ),dPQ(a,b)=1S∑s=0S−1ds(a(s),b(s))d_{\text{PQ-Chamfer}}(q, D) = \text{Chamfer}\left(\{q\}, D; \; d_{\text{PQ}}\right), \quad d_{\text{PQ}}(a, b) = \frac{1}{S} \sum_{s=0}^{S-1} d_s(a^{(s)}, b^{(s)})

VT-Aligned (Chamfer-first). VT-Aligned reverses this order: for each subspace ss independently, it first computes the sentence-level Chamfer distance (in dsd_s dimensions), then averages across subspaces: dVT(q,D)=1S∑s=0S−1[minj∈Dds(q(s),sj(s))+1|D|∑j∈Dds(sj(s),q(s))]d_{\text{VT}}(q, D) = \frac{1}{S} \sum_{s=0}^{S-1} \left[\min_{j \in D} d_s(q^{(s)}, s_j^{(s)}) + \frac{1}{|D|} \sum_{j \in D} d_s(s_j^{(s)}, q^{(s)}) \right] where ds(a,b)=1−cos(a(s),b(s))d_s(a, b) = 1 - \cos(a^{(s)}, b^{(s)}) is the cosine distance restricted to the ss-th 64-dimensional subspace.

Physical interpretation. Each 64-dimensional subspace acts as an independent "virtual token"—a semantic dimension that captures one facet of the full embedding. The query's single 4096d vector is effectively decomposed into 64 independent semantic probes, each searching for its best sentence match within its own subspace. This increases the effective query representation density from 1 point to 64 points without any model re-inference or additional data. The key insight is that because Qwen3-8B distributes semantic information uniformly across embedding dimensions (confirmed by ablation: dimension permutation has negligible effect), each subspace constitutes a legitimate independent semantic channel.

VT-Aligned vs. VT-Unaligned. We also evaluate an unaligned variant that flattens all subspaces into a single pool of 64d virtual tokens and allows cross-subspace matching (subspace ii of the query can match subspace j≠ij \neq i of a document sentence). VT-Aligned consistently outperforms VT-Unaligned (+2.5% vs +0.9% over PQ-Chamfer baseline on NFCorpus), confirming that cross-subspace matching introduces noise rather than additional signal.

3.8.1 Adjoint State Prefetch (V7.1)

The adjoint state prefetch method addresses the candidate ceiling limitation of reranking: a reranker can only reorder documents already present in the initial retrieval set. V7.1 identifies high-pressure boundary nodes—documents with high post-PDE concentration but positioned at the semantic frontier of the candidate set—and uses them as probes to expand the candidate pool.

Boundary node selection. After the initial PDE pass, boundary nodes are selected via the adjoint flux criterion: $i^* = \arg\max_{i \in \partial V} C_i \cdot \max\left(0, \sum_{j \in \mathcal{N}(i)} U_{ij} W_{ij} C_j\right)$ where ∂V denotes nodes with graph degree below the median. This selects nodes that simultaneously have high relevance concentration and net inward advective flux—indicating that the PDE dynamics are "pulling" additional relevance from beyond the candidate boundary.

Probe construction and retrieval. For each selected boundary node d_i, a probe vector is constructed via cosine interpolation: $\text{probe}_i = 0.7 \hat{q} + 0.3 \hat{d}_i$ where $\hat{q}$ and $\hat{d}_i$ are unit-normalized. The probe is used to retrieve additional candidates from the full corpus via cosine similarity, expanding the candidate pool before a second PDE pass.

3.9 Token-Level Point Clouds (V11)

The sentence-level point cloud formulation of Shape-CFD (Section 3.6) represents each document as a collection of 5–10 sentence embeddings. While this captures inter-sentence semantic structure, the granularity remains coarse: each sentence embedding is a pooled summary of dozens of tokens, compressing fine-grained lexical and semantic distinctions into a single vector. More critically, single-sentence queries degenerate to a single point in embedding space, eliminating any geometric structure on the query side—a limitation we term "single-point query degeneracy."

Token-level representation. We propose extracting per-token last hidden states from the embedding model (rather than the pooled [CLS] or mean-pooled output), yielding dense point clouds at the token level: $d_i \to \mathcal{T}_i = \{t_{i,1}, t_{i,2}, \ldots, t_{i,M_i}\} \subset \mathbb{R}^d$ where M_i is the number of tokens in document d_i (typically 100–400+ tokens), and each $t_{i,k} \in \mathbb{R}^{4096}$ is the last hidden state of the k-th token. Similarly, the query is expanded into a token-level point cloud: $q \to \mathcal{T}_q = \{t_{q,1}, t_{q,2}, \ldots, t_{q,L}\} \subset \mathbb{R}^d$ where L is the number of query tokens (typically $6 for short queries). This resolves the single-point degeneracy: even a 3-word query produces 6 token vectors, each capturing a distinct semantic facet.

Token PQ-Chamfer distance. We extend the PQ-Chamfer decomposition (Section 4.3.5) to token-level point clouds. The embedding space is partitioned

into S=64S = 64 contiguous 64-dimensional subspaces, and the Token PQ-Chamfer distance is defined as:

dTC(Q,D)=1S∑s=1S[1|Q|∑q∈Qmind∈Dcs(q,d)+1|D|∑d∈Dminq∈Qcs(q,d)]d_{TC}(Q, D) = \frac{1}{S} \sum_{s=1}^{S} \left[\frac{1}{|Q|} \sum_{q \in Q} \min_{d \in D} c_s(q, d) + \frac{1}{|D|} \sum_{d \in D} \min_{q \in Q} c_s(q, d) \right]

where cs(q,d)=1−cos(qs,ds)c_s(q, d) = 1 - \cos(q_s, d_s) is the cosine distance within subspace ss, and qs,dsq_s, d_s denote the ss-th 64-dimensional slice of the respective token vectors.

Relationship to cross-encoders. The Token PQ-Chamfer distance can be interpreted as a zero-parameter hard cross-attention mechanism. Each query token attends to its single most similar document token (and vice versa) within each subspace independently—this is the min\min operator in the Chamfer formulation. By contrast, a cross-encoder computes soft attention with learned, parameterized weights over all query-document token pairs. The key distinction is that Token PQ-Chamfer requires no training, no parameters, and no joint query-document encoding—it operates purely on pre-computed token embeddings. From an Unbalanced Optimal Transport (UOT) perspective, the token-level Chamfer corresponds to a hard assignment transport plan where each token is matched to its nearest counterpart with zero regularization, as opposed to the entropy-regularized Sinkhorn transport that underpins soft attention.

V12 falsification and query-side resolution. An earlier variant (V12) attempted to apply token-level matching with sentence-level PDE graph structure, using token Chamfer only for the initial scoring while retaining the sentence-level KNN graph for diffusion. This produced NDCG@10 = 0.3180—inferior to pure token Chamfer (0.3214)—because the sentence-level graph topology is mismatched with the token-level concentration field. The semantic neighborhoods defined by sentence centroids do not align with the fine-grained token-level distance manifold, causing the PDE to diffuse concentration along edges that are irrelevant at the token scale. This falsification motivates the fusion approach (Section 3.9.2) rather than hierarchical integration.

3.9.1 Two-Stage Token Retrieval

Computing full Token PQ-Chamfer distances over the entire corpus is expensive: with $$356 tokens per document and 3,633 documents in NFCorpus, a full scan requires O(N·M·L·S)O(N \cdot M \cdot L \cdot S) distance computations, resulting in $$305 ms per query. We introduce a two-stage acceleration strategy that achieves lossless accuracy with a 13x speedup.

Stage 1: Token centroid coarse filtering. For each document, we precompute the mean of all token vectors (the "token centroid"): t¯i=1Mi∑k=1Mitiˉ,k\bar{t}_i = \frac{1}{M_i} \sum_{k=1}^{M_i} t_{i,k} The query centroid t¯q\bar{t}_q is computed analogously. Stage 1 ranks all corpus documents by cosine similarity between token centroids and selects the top-K1K_1 candidates (we use K1=100K_1 = 100).

Stage 2: Full token PQ-Chamfer re-ranking. The top-K1K_1 candidates from Stage 1 are re-ranked using the full Token PQ-Chamfer distance (computing the complete token-to-token distance matrix within each subspace). The final output is the top-K2K_2 documents (we use K2=55K_2 = 55, consistent with the sentence-level pipeline).

Lossless acceleration. On NFCorpus (323 queries), the two-stage pipeline achieves NDCG@10 = 0.3220 versus 0.3214 for full corpus scan—a marginal improvement of +0.2%, attributable to the centroid pre-filter excluding distant documents whose token-level noise would otherwise perturb the ranking. Latency drops from 305 ms to 23 ms (13.3x speedup), making the token-level pipeline practical for real-time deployment.

3.9.2 Fusion: PDE as Orthogonal Signal Source

A natural question is whether PDE-based reranking can further improve token Chamfer rankings, as it did for sentence-level Shape-CFD. We investigate three integration strategies and discover that the correct role of PDE at the token level is orthogonal signal fusion, not same-perspective smoothing.

Negative result: Token + sentence PDE (V12). Applying sentence-level PDE reranking to token Chamfer initial scores yields NDCG@10 = 0.3180, worse than pure token Chamfer (0.3214). The cause is a fundamental semantic space mismatch: the sentence-level KNN graph captures coarse inter-document neighborhoods, but the token-level initial concentrations encode fine-grained token matches. Diffusion along sentence-level edges smooths away the discriminative token-level signal rather than enhancing it.

Negative result: Sampled graph PDE. Constructing a KNN graph from sampled token subsets (to approximate token-level topology) yields 0.3180, confirming that the graph-concentration mismatch is not merely a resolution issue but a structural incompatibility between graph-based propagation and token-level scoring.

Positive result: Score-level fusion. Instead of applying PDE on top of token scores (sequential pipeline), we compute PDE scores and token Chamfer scores

independently and combine them via linear fusion: sfusion(di)=λ·sTC(di)+(1−λ)·sPDE(di)s_{\text{fusion}}(d_i) = \lambda \cdot s_{TC}(d_i) + (1-\lambda) \cdot s_{PDE}(d_i) where sTCs_{TC} is the normalized Token PQ-Chamfer score (from the two-stage pipeline) and sPDEs_{PDE} is the normalized Shape-CFD v10 Rust pipeline score. The two scores are independently min-max normalized to [0,1][0, 1] before fusion.

Optimal fusion weight. A sweep over λ∈{0.3,0.5,0.7,0.9}\lambda \in \{0.3, 0.5, 0.7, 0.9\} reveals the optimum at λ=0.7\lambda = 0.7:

λ\lambda	NDCG@10	Note
1.0 (pure Token)	0.3220	Two-stage token Chamfer alone
0.9	0.3224	Marginal PDE contribution
0.7	**0.3232**	**Optimal fusion**
0.5	0.3210	Equal weight—PDE noise begins to dominate
0.3	0.3175	PDE-dominant—degrades
0.0 (pure PDE)	0.2852	Rust v10 pipeline alone

The optimal λ=0.7\lambda = 0.7 assigns 70% weight to the token Chamfer signal (local, fine-grained nearest-neighbor matching) and 30% to the PDE signal (global, graph-based semantic propagation). This reflects the complementary nature of the two signals: token Chamfer excels at identifying documents with precise lexical/semantic overlap, while PDE excels at promoting documents that are globally well-connected in the semantic graph but may lack a direct token-level match with the query.

Why PDE helps as fusion but hurts as pipeline. When PDE is applied sequentially (as a reranker on top of token scores), it operates within the token-scored ranking and can only redistribute scores among documents already ranked by token similarity. This smoothing destroys the fine-grained token-level discriminations. When PDE is applied in parallel (as an independent scorer), it provides a different view of document relevance—one based on global semantic topology rather than local token matching. The fusion then selects documents that are highly ranked by both criteria, achieving a robustness that neither signal provides alone.

3.10 Dual Physical Indicators

We introduce two complementary dimensionless numbers to characterize the algorithm’s operating regime:

Dynamic Reynolds number (runtime indicator): Re=∑i=1N|Fiadv|∑i=1N|Fidiff| \text{Re} = \frac{\sum_{i=1}^{N} |F_i^{\text{adv}}|}{\sum_{i=1}^{N} |F_i^{\text{diff}}|} This measures the ratio of advective to diffusive activity (taking absolute values at each node to avoid the trivial cancellation inherent in symmetric graphs).

Structural Péclet number (static indicator): Pe=|uij|¯D=1|ℰ|∑(i,j)∈ℰ|uij|D\text{Pe} = \frac{\overline{|u_{ij}|}}{D} = \frac{\frac{1}{|\mathcal{E}|} \sum_{(i,j) \in \mathcal{E}} |u_{ij}|}{D} This is computed once at graph construction time, independent of the concentration field CC, and captures the intrinsic balance between advective and diffusive capacity of the graph structure.

---

4 Experiments

4.1 Datasets

We evaluate Shape-CFD on four datasets spanning two languages and three domains:

BEIR Benchmarks. We use three subsets from the BEIR benchmark suite [25], widely adopted for evaluating zero-shot retrieval:

Dataset	Domain	Queries	Corpus Size	Avg. Doc Length
SciFact	Scientific claims	300	5,183	Short abstracts
NFCorpus	Biomedical	323	3,633	Mixed lengths
FiQA	Financial QA	648	57,638	Variable

LawVein Legal Corpus. A proprietary Chinese legal knowledge base consisting of 23,701 statute chunks extracted from 24 Chinese laws, encoded with Qwen3-Embedding-8B (d=4096d = 4096) and indexed using USearch HNSW (M=32M = 32, ef=128\text{ef} = 128). This dataset serves as our primary development testbed and provides a challenging scenario with highly homogeneous intra-domain embeddings.

4.2 Baselines and Evaluation Protocol

Retrieval methods compared:

- **Cosine**: Direct HNSW retrieval with cosine similarity ranking (the standard baseline).
- **BM25**: Sparse lexical matching via Okapi BM25.

- **AD-Rank v2**: Our point-based convection-diffusion algorithm (without the shape extension), using optimal parameters D=0.15D = 0.15, ustrength=0.1u_{\text{strength}} = 0.1, k=3k = 3.
- **Shape-CFD**: The full shape-aware extension with Chamfer distance, JL projection (m=128m = 128), and sentence-level point clouds.

Strategy variants (for the LawVein ablation study):

Variant	Description	Key Parameter
MDA	Multi-Directional Advection: 8 random projection subspaces, each with independent graph and advection	M=8M = 8, k=32k = 32
BAA	Block-Adaptive Advection: 8 contiguous dimension blocks	B=8B = 8
v2+PCA	Local PCA dimensionality reduction for advection direction	pcaDim=32\text{pcaDim} = 32
v2+Temperature	Nonlinear temperature scaling of advection coefficients	α=0.5\alpha = 0.5
v2+Adaptive	Fiedler-value adaptive DD + cross-domain adaptive uu	Auto

Evaluation metrics. For BEIR benchmarks, we report NDCG@10, MRR@10, and Recall@10 using the official evaluation scripts. For the LawVein dataset (which lacks ground-truth relevance labels), we employ **DeepSeek V3.2 blind evaluation**: an LLM judge independently scores anonymized A/B result pairs on a 1–5 scale for relevance, coverage, and ranking rationality, with random assignment to prevent position bias.

4.3 Main Results

BEIR Benchmark Results

Dataset	Cosine	BM25	AD-Rank v2	Shape-CFD	V7.1 Adjoint	VT-V7 (JS)	Rust Pipeline	Token 2-Stage	Fusion	Δ\Delta best vs Cosine
SciFact	0.4701	—	0.4753	0.4820	—	—	—	—	—	+2.5%

Dataset	Cosine	BM25	AD-Rank v2	Shape-CFD	V7.1 Adjoint	VT-V7 (JS)	Rust Pipeline	Token 2-Stage	Fusion	Δ\Delta best vs Cosine
NFCorpus	0.2195	—	0.2398	0.2684	0.2802	0.2844	0.2852	0.3220	**0.3232**	**+47.2%**
FiQA	0.1683	—	0.1683	0.1683	—	—	—	—	—	0%

V7.1 Adjoint denotes the adjoint state prefetch method (Section 3.8.1), which uses PDE concentration gradients to identify high-pressure boundary nodes and expand the candidate pool via speculative probing, achieving NDCG@10 = 0.2802 (+27.6% vs cosine) on NFCorpus. **VT-V7 (JS)** denotes the JavaScript stacked configuration of VT-Aligned virtual token decomposition (Section 3.7.1) and V7.1 adjoint prefetch, achieving NDCG@10 = 0.2844 (+29.6% vs cosine) on NFCorpus. **Rust Pipeline** denotes the calibrated Rust full-pipeline engine (Section 4.6) with 6 bug fixes and centroid-based cosine pre-filter (top_n=55), achieving NDCG@10 = 0.2852 (+29.9% vs cosine). **Token 2-Stage** denotes the two-stage token-level pipeline (Section 3.9.1) with centroid coarse filter (top-100) and full token PQ-Chamfer re-ranking, achieving NDCG@10 = 0.3220 (+46.7% vs cosine) in 23 ms. **Fusion** denotes the score-level fusion (Section 3.9.2) of token 2-stage and Rust PDE pipeline with λ=0.7\lambda = 0.7, achieving the overall best NDCG@10 = **0.3232** (+47.2% vs cosine) at 26 ms latency.

On SciFact and NFCorpus, Shape-CFD substantially outperforms both the cosine baseline and the point-based AD-Rank v2. The NFCorpus improvement of +21.5% (Shape-CFD alone) is particularly notable, as NFCorpus contains biomedical documents with rich multi-sentence structure that benefits from the Chamfer distance-based graph construction. The subsequent VT-Aligned and adjoint prefetch innovations push the total gain to +29.9% (Rust calibrated pipeline). The V11 token-level extension represents a paradigm-level leap: by operating on per-token hidden states rather than pooled sentence embeddings, Token PQ-Chamfer achieves +46.4% over cosine via full scan, and the two-stage centroid acceleration reaches +46.7% in only 23 ms. The final fusion with PDE as an orthogonal signal source achieves +47.2%—nearly half again the improvement of the sentence-level pipeline—at 26 ms latency on CPU with zero cross-encoder inference.

The FiQA result (Δ=0\Delta = 0) is informative: FiQA’s financial Q&A documents tend to be short, single-sentence answers that yield singleton point clouds. In this regime, Shape-CFD degrades exactly to cosine similarity (as expected), confirming that the shape extension contributes only when multi-sentence structure is present—it does not introduce spurious effects on incompatible document types.

V11 Complete Ablation (NFCorpus, 323 queries)

The following table presents the complete ablation across all 10 evaluated methods on NFCorpus, ordered by NDCG@10:

Method	NFCorpus NDCG@10	vs Cosine	Latency
Cosine baseline	0.2195	—	20 ms
BM25	0.2573	+17.2%	—
AD-Rank v2	0.2398	+9.2%	—
Shape-CFD (sentence)	0.2684	+22.3%	—
V7.1 Adjoint	0.2802	+27.6%	—
VT-V7 (JS)	0.2844	+29.6%	—
Rust Pipeline v10	0.2852	+29.9%	26 ms
Token Chamfer (full scan)	0.3214	+46.4%	305 ms
Token 2-Stage (coarse=100)	0.3220	+46.7%	23 ms
Fusion (λ\lambda=0.7)	**0.3232**	**+47.2%**	**26 ms**

Key observations from the ablation:

1. **Token-level is a paradigm shift.** The jump from Rust Pipeline v10 (0.2852) to Token Chamfer full scan (0.3214) is +12.7%—larger than the entire cumulative gain from AD-Rank v2 through VT-V7 stacking. This confirms that per-token hidden states contain substantially more discriminative information than pooled sentence embeddings.

2. **Two-stage is lossless.** Token 2-Stage (0.3220) marginally exceeds full scan (0.3214), demonstrating that centroid pre-filtering is not merely an approximation but an effective noise filter that excludes distant documents whose token-level contributions are purely noise.

3. **PDE contributes orthogonally.** The fusion gain (+0.0012, from 0.3220 to 0.3232) is modest but consistent, and arises from the complementary nature of local token matching and global graph propagation. Notably, this gain is achieved at essentially zero additional latency (PDE scores are precomputed by the Rust pipeline).

LawVein Blind Evaluation Results

AD-Rank v2 vs. Cosine (10 queries, DeepSeek V3.2 blind judge):

Method	Mean Score	Win	Loss	Draw	Win Rate
AD-Rank v2	3.4/5	7	2	1	**70%**
Cosine	2.8/5	2	7	1	20%

Shape-CFD vs. Cosine (10 queries, independent evaluation):

Method	Mean Score	Win	Loss	Win Rate
Shape-CFD	3.4/5	**8**	2	**80%**
Cosine	2.6/5	2	8	20%

Shape-CFD achieves an 80% win rate against cosine (8:2), surpassing AD-Rank v2’s 70% win rate (7:2:1). The improvement is most pronounced on cross-domain queries (e.g.,“divorce property division”spanning civil law and procedural law), where the directional advection and multi-sentence matching provide the greatest benefit.

Full Strategy Cross-Evaluation (LawVein)

To validate AD-Rank v2 as the optimal point-based configuration, we conducted a comprehensive 7-method tournament with 10×6=6010 \times 6 = 60 blind comparisons:

Rank	Method	Blind Score	Pe	Latency	Convergence	W/L
🥇	**AD-Rank v2**	**3.4/5**	0.065	14.8 ms	**100%**	**7/2**
🥈	BAA	3.3/5	0.226	7 ms	20%	7/3
🥉	v2+PCA	3.2/5	0.112	33 ms	90%	6/4
4	v2+Adaptive	3.0/5	0.122	3.3 ms	100%	5/4
5	v2+Temperature	2.9/5	0.604	3.5 ms	10%	5/5
6	MDA	2.7/5	0.351	85.8 ms	0%	3/7

Key findings: (1) AD-Rank v2 achieves the highest blind evaluation score with 100% convergence. (2) Higher Péclet numbers do not correlate with better ranking quality—the methods with the strongest advection (MDA: Pe=0.351, Temperature: Pe=0.604) perform worst, supporting our hypothesis that weak advection is optimal. (3) Convergence rate is a hard constraint: methods with convergence rates below 50% (BAA, Temperature, MDA) produce inconsistent results.

4.3.3 V5 Reaction-Diffusion Ablation (v3 supplement)

In Section 6 of v1, we proposed augmenting the convection-diffusion equation with an Allen-Cahn bistable reaction term R(C)=γC(1−C)(C−θ)R(C) = \gamma\, C(1-C)(C-\theta) to produce genuinely heterogeneous stationary distributions. We report systematic experiments testing this proposal. All experiments in Sections 4.3.3–4.3.4 use a unified evaluation framework operating on the native d=4096d = 4096 embedding space (without JL projection), with the Chamfer-distance graph topology and N=30N = 30, k=3k = 3.

V5.0: Unnormalized Allen-Cahn. We introduce the reaction term with γ=0.5\gamma = 0.5 and threshold θ\theta set to the mean of the initial concentration field. On NFCorpus, this yields NDCG@10 = 0.2467 (−1.6%-1.6\% vs. Shape-CFD v4 baseline of 0.2508), confirming that the naive addition of reaction dynamics degrades performance. The root cause: when C0∈[0.1,0.2]C_0 \in [0.1, 0.2] (a typical range for cosine-transformed scores), the reaction term evaluates to R(C)≈0.004R(C) \approx 0.004, which after multiplication by Δt=0.05\Delta t = 0.05 contributes only ∼0.0002\sim 0.0002 per step—completely overwhelmed by the diffusion flux.

V5.1: Normalized Allen-Cahn. Following diagnostic analysis, we attempted to amplify the reaction dynamics by applying Min-Max normalization to map C0C_0 to [0,1][0,1] before applying R(C)R(C) with γ=5.0\gamma = 5.0. This produced a catastrophic −12.9%-12.9\% regression on NFCorpus (NDCG@10 = 0.2184). The failure mechanism: Min-Max normalization stretches a natural range of Δ≈0.1\Delta \approx 0.1 to the full [0,1][0,1] interval—an effective spatial dilation of 10×10\times. The combination with γ=5.0\gamma = 5.0 yields an effective reaction rate of γeff≈500\gamma_{\text{eff}} \approx 500 in the original coordinate system, triggering violent phase separation that destroys the carefully calibrated ranking signal. This result empirically confirms that γ\gamma must be scaled inversely with the square of the score range.

V5.2: Domain-adapted Allen-Cahn. We implement the corrected formulation: γeff=γbase(Cmax−Cmin+ϵ)2,R(C)=γeff(C−Cmin)(Cmax−C)(C−θOtsu) \gamma_{\text{eff}} = \frac{\gamma_{\text{base}}}{(C_{\max} - C_{\min} + \epsilon)^2}, \quad R(C) = \gamma_{\text{eff}} (C - C_{\min})(C_{\max} - C)(C - \theta_{\text{Otsu}})

where θOtsu\theta_{\text{Otsu}} is obtained via Otsu’s method on the initial concentration histogram. We additionally implement Zelnik-Manor local bandwidth scaling: Wij=Wij0·exp(−∥vi−vj∥2/(σiσj))W_{ij} = W_{ij}^0 \cdot \exp(-\|v_i - v_j\|^2 / (\sigma_i \sigma_j)), where σi\sigma_i is the distance from node ii to its kk-th nearest neighbor.

Configuration	NFCorpus NDCG@10	Δ\Delta vs V4
Shape-CFD v4 (baseline)	0.2508	—
V5.0: γ=0.5\gamma=0.5, unnormalized	0.2467	−1.6%-1.6\%
V5.1: γ=5.0\gamma=5.0, normalized	0.2184	−12.9%-12.9\%
V5.2: γbase=0.5\gamma_{\text{base}}=0.5, domain-adapted	0.2501	−0.3%-0.3\%
V5.2: γbase=1.0\gamma_{\text{base}}=1.0, domain-adapted	0.2501	−0.3%-0.3\%
V5.2: γbase=2.0\gamma_{\text{base}}=2.0, domain-adapted	0.2501	−0.3%-0.3\%

The domain-adapted formulation (V5.2) successfully prevents the catastrophic phase transition of V5.1, recovering to within 0.3%0.3\% of the baseline. However, three values of γbase\gamma_{\text{base}} (0.5, 1.0, 2.0) produce identical results (0.2501), demonstrating that the reaction term’s contribution is negligible—the system is **topologically locked** to a Dirichlet energy minimum where the small-sample KNN graph (N=30N=30, diameter ≤4\leq 4) allows diffusion to equilibrate instantaneously, rendering any perturbation from the reaction term invisible to the NDCG ranking metric.

Conclusion. In the short-query, narrow score-range regime (ΔC0≈0.1\Delta C_0 \approx 0.1), the Allen-Cahn reaction term is numerically dormant—not due to implementation error, but as a physical necessity. The reaction force scales as O(Δ3)O(\Delta^3) while diffusion scales as O(Δ)O(\Delta); when Δ≪1\Delta \ll 1, diffusion dominates by construction. This result validates the v1 observation (Section 5.1) that the optimal operating regime is diffusion-dominated, and precisely delineates the boundary conditions under which nonlinear dynamics become physically meaningful: only when the initial score range is sufficiently large (ΔC0≫0.1\Delta C_0 \gg 0.1), as expected in long-document retrieval scenarios with sparse evidence.

4.3.4 Mixed Initial Field Ablation (v3 supplement)

The v1 Shape-CFD formulation uses the symmetric Chamfer distance to compute both graph weights and initial concentrations. For single-vector queries, the query-document Chamfer distance dChamfer(q,𝒫i)d_{\text{Chamfer}}(q, \mathcal{P}_i) decomposes into two components: dChamfer(q,𝒫i)=mins∈𝒫idcos(q,s)__MaxSim (forward)+1|𝒫i| ∑s∈𝒫idcos(s,q)__MeanSim (reverse)d_{\text{Chamfer}}(q, \mathcal{P}_i) =

$$\underbrace{\min_{s \in \mathcal{P}_i} d_{\cos}(q, s)}_{\text{MaxSim (forward)}} + \underbrace{\frac{1}{|\mathcal{P}_i|} \sum_{s \in \mathcal{P}_i} d_{\cos}(s, q)}_{\text{MeanSim (reverse)}}$$

The first term captures the best local match (the most relevant sentence), while the second penalizes documents with many unmatched sentences. In v1, these components are implicitly weighted equally ($\alpha = 0.5$). We investigate whether decoupling this balance can improve performance.

Mixed initial field. We replace the Chamfer-based initial concentration with an explicit mixture: $C_0(i) = \alpha \cdot \text{MaxSim}(q, \mathcal{P}_i) + (1-\alpha) \cdot \text{MeanSim}(q, \mathcal{P}_i)$ where $\text{MaxSim}(q, \mathcal{P}_i) = \max_{s \in \mathcal{P}_i} \cos(q, s)$ and $\text{MeanSim}(q, \mathcal{P}_i) = \frac{1}{|\mathcal{P}_i|} \sum_{s \in \mathcal{P}_i} \cos(q, s)$. The graph topology remains Chamfer-based (preserving the high-fidelity Riemannian manifold), and the PDE solver uses pure diffusion (V4 mode, no reaction term).

Results across two BEIR datasets:

α	NFCorpus	Δ	SciFact	Δ
0.5 ($\approx$ Chamfer)	0.2539	+1.2%	0.4431	−1.5%
0.6	0.2567	+2.3%	0.4434	−1.5%
0.7	**0.2588**	**+3.2%**	0.4413	−1.9%
0.8	0.2578	+2.8%	0.4411	−2.0%
0.9	0.2546	+1.5%	0.4388	−2.5%
1.0 (pure MaxSim)	0.2535	+1.1%	0.4319	−4.0%

Baselines: NFCorpus Shape-CFD v4 = 0.2508; SciFact Shape-CFD v4 = 0.4500.

The results reveal a striking dataset-dependent divergence:

- **NFCorpus** (heterogeneous, multi-sentence biomedical documents): The α response follows an inverted-U curve with a peak at $\alpha = 0.7$ (+3.2%). MaxSim-dominated initialization penetrates the signal dilution caused by long, noisy documents, while the 30% MeanSim component provides a "macroscopic gravity" that prevents false-positive promotion.
- **SciFact** (homogeneous, short scientific abstracts): The response is monotonically decreasing—the higher α is, the worse the performance, with pure MaxSim ($\alpha = 1.0$) degrading by

−4.0%-4.0\%. In 3–5 sentence documents where every sentence is signal, MaxSim ≈\approx MeanSim, so decoupling adds no information but introduces extreme-value sampling noise from picking a single sentence.

This asymmetry is analyzed theoretically in Section 5.4.

4.4 Ablation Study

Impact of the Shape Extension

We isolate the contribution of the shape (point cloud) representation by comparing three configurations on the BEIR datasets:

Configuration	SciFact	NFCorpus	FiQA	Description
Cosine (no PDE)	0.4701	0.2209	0.1683	No graph dynamics
AD-Rank v2 (point PDE)	0.4753	0.2398	0.1683	Graph dynamics, point embeddings
Shape-CFD (shape PDE)	**0.4820**	**0.2684**	0.1683	Graph dynamics, point clouds

The shape extension provides the marginal gain over AD-Rank v2: +1.4% on SciFact, +11.9% on NFCorpus. The FiQA null result confirms the hypothesis: the shape extension is selectively beneficial, activating only when documents contain multi-sentence structure. This dataset-adaptive behavior is a strength, not a limitation—it ensures no degradation on short-document corpora.

Péclet Number Sensitivity

We vary the advection strength ustrength∈{0,0.05,0.1,0.2,0.5,1.0} u_{\text{strength}} \in \{0, 0.05, 0.1, 0.2, 0.5, 1.0\} while fixing D=0.15D = 0.15 on the LawVein dataset:

ustrengthu_{\text{strength}}	Pe	Convergence	Blind Score	Ranking
0.00	0.000	✓	3.1	Pure diffusion
0.05	0.033	✓	3.2	Under-advected
0.10	**0.065**	✓	**3.4**	**Optimal**
0.20	0.130	✓	3.1	Over-advected
0.50	0.330	✗	2.8	Non-convergent
1.00	0.650	✗	2.3	Divergent

The optimal operating point is at Pe≈0.065\text{Pe} \approx 0.065, in the weakly advective regime. This result is robust across queries: the ustrength=0.1u_{\text{strength}} = 0.1 configuration achieves 100% convergence while consistently outperforming both pure diffusion (u=0u = 0) and stronger advection variants.

4.5 Latency Analysis

End-to-end latency measurements on a single-core AMD EPYC (Zen 4, 4.2 GHz):

Component	AD-Rank v2	Shape-CFD
HNSW retrieval (30 candidates)	1.2 ms	1.2 ms
Sentence embedding (API)	—	1,815 ms (uncached)
JL projection + caching	—	0.3 ms (cached)
Graph construction	0.4 ms	2.1 ms
CFD solving (30-50 iterations)	1.0 ms	1.5 ms
Total (cached)	**2.6 ms**	**5.1 ms**
Total (with reranking overhead only)	**1.4 ms**	**3.9 ms**

Key observations: (1) The raw PDE solving time is negligible (⩽1.5\leq 1.5 ms), confirming that the algorithm’s computational core is extremely efficient for the problem scale (N=30N = 30). (2) The primary bottleneck for Shape-CFD is the **one-time** sentence embedding cost (1,815 ms per document via API); once cached in SQLite, subsequent queries incur only the 0.3 ms projection lookup. (3) Both methods are approximately 100×100\!\times–200×200\!\times faster than cross-encoder rerankers (500+ ms for N=30N = 30).

Optimization trajectory. Over the course of development, we achieved a cumulative 300×300\!\times speedup through five stages of optimization:

Stage	Latency	Optimization	Factor
v0 (baseline)	450 ms	No optimization	1×
v1	120 ms	Pre-compute similarity matrix	3.75×
v2	14.8 ms	KNN sparse graph (k=3k=3)	8.1×
v3	3.3 ms	Adaptive early stopping	4.5×
v4 (Shape-CFD)	22 ms	+Shape with JL (m=128m=128) + cache	—
v4.1 (VT-V7 Rust)	**38 ms**	Full-pipeline Rust: fullscan + VT + PDE + adjoint	—

4.6 Rust Full-Pipeline Engine (v4 supplement)

The VT-V7 configuration (VT-Aligned + V7.1 adjoint prefetch) constitutes the final production pipeline. To eliminate all legacy dependencies, we implement the complete retrieval-to-reranking pipeline as a native Rust module (`law-vexus`), with the following architectural decisions:

Complete departure from HNSW and cosine recall. The Rust engine fully abandons HNSW approximate nearest-neighbor search and cosine-similarity-based first-stage recall. Instead, it performs **fullscan** (brute-force exhaustive scan) over the entire corpus using VT-Aligned distance as the primary ranking metric, followed by PDE reranking. This design is motivated by the observation that HNSW’s approximate recall introduces a candidate ceiling that the adjoint prefetch must compensate for; eliminating HNSW removes this ceiling entirely.

Module architecture. The engine comprises five core modules (total ~2,148 lines of Rust):

Module	Lines	Responsibility
`cloud_store.rs`	406	Point cloud storage, sentence embedding management, norm pre-caching
`fullscan.rs`	517	Exhaustive corpus scan with VT-Aligned distance, rayon parallel iteration
`pq_chamfer.rs`	382	PQ subspace decomposition, 64x64 structured Chamfer distance
`vt_distance.rs`	605	VT-Aligned and VT-Unaligned distance kernels, 4-way unrolled SIMD inner loops
`pde.rs`	238	Graph construction, CFL-adaptive time stepping, conservative upwind PDE solver

Performance optimizations:

1. **Rayon parallel fullscan.** The corpus scan is parallelized across all available CPU cores via `rayon::par_iter`, with each thread computing VT-Aligned distances for a partition of the corpus independently.

2. **4-way loop unrolling with SIMD-friendly layout.** The innermost distance computation loop is manually unrolled 4x, with memory layout aligned to 64-byte cache lines to enable auto-vectorization by the LLVM backend.

3. **Norm pre-caching.** All document sentence norms are precomputed and stored in `cloud_store.rs` at index time, eliminating redundant ∥v∥\|v\| computations during online distance evaluation.

4. **Struct-of-Arrays (SoA) memory layout.** Sentence embeddings are stored in SoA format (one contiguous array per subspace) rather than Array-of-Structs, maximizing cache locality for the subspace-sequential VT-Aligned computation.

End-to-end latency: 38 ms per query on commodity CPU hardware (AMD EPYC, single socket), representing a 130x speedup over the JavaScript reference implementation. This latency includes fullscan, VT-Aligned graph construction, PDE solving, and adjoint prefetch—the complete retrieval-to-reranking pipeline.

Test coverage. The module includes 45 unit tests covering distance computation correctness, PDE conservation invariants, CFL bound verification, and end-to-end ranking regression tests.

4.7 V11 Token-Level Implementation

The token-level pipeline (V11) requires a complete re-extraction of document and query representations at the token granularity, along with new storage and computation infrastructure.

Token extraction. We use `llama.cpp` with `--pooling none` to extract per-token last hidden states from a Q4_K_M quantized Qwen3-Embedding-8B model. With pooling disabled, the model outputs one 4096-dimensional vector per input token (rather than a single pooled vector). For NFCorpus, this yields 880,791 token vectors across 3,633 documents (average $$356 tokens per document, including BOS/EOS special tokens) and 1,976 query token vectors across 323 queries (average $$6 tokens per query).

Storage. Token embeddings are stored in SQLite using f32 native-endian BLOBs, with a schema compatible with the existing sentence-level point cloud database:

```
CREATE TABLE token_clouds (
  doc_id TEXT PRIMARY KEY,
  n_tokens INTEGER,
  embeddings BLOB  -- M_i * 4096 * 4 bytes, f32 native endian
);
CREATE TABLE query_token_clouds (
  query_id TEXT PRIMARY KEY,
  n_tokens INTEGER,
  embeddings BLOB
);
```

The full token cloud database for NFCorpus is approximately 14 GB (880,791 tokens ×\times 4096 dimensions ×\times 4 bytes).

Rust engine modules. The V11 pipeline adds two core modules to the `law-vexus` Rust engine:

Module	Responsibility
`token_chamfer.rs`	Token PQ-Chamfer distance computation with distance matrix reuse, prenorm optimization, and 64-subspace decomposition
`extract_tokens.rs`	Binary tool for batch token extraction from llama.cpp output, SQLite serialization, and centroid precomputation

Distance matrix reuse. For the two-stage pipeline, the token-to-token cosine distance matrix within each subspace is computed once and reused for both the forward (Q→DQ \to D) and reverse (D→QD \to Q) Chamfer terms. With prenorm (all token vectors are L2L_2-normalized at load time), the cosine distance reduces to 1−q·d1 - q \cdot d, enabling efficient BLAS-level matrix multiplication for batch distance computation.

Centroid precomputation. Document token centroids (the mean of all token vectors) are precomputed at indexing time and stored alongside the full token clouds. This enables the Stage 1 coarse filter to operate as a simple cosine similarity lookup over centroid vectors, with zero additional computation at query time.

4.8 Case Study: Cross-Domain Legal Query

We present a representative query to illustrate Shape-CFD’s qualitative advantage:

Query: “离婚时一方隐藏财产怎么处理” (How to handle hidden property during divorce)

This query spans two legal domains: (1) marriage law (Article 1092 of Civil Code on concealment of marital property) and (2) civil procedure law (evidence preservation provisions). A pure cosine ranker retrieves documents from both domains but fails to recognize their complementary relationship.

Cosine ranking: The top-3 results correctly include Article 1092 but rank it at position 3, behind two marginally relevant provisions from the general section of marriage law.

Shape-CFD ranking: By constructing a Chamfer-distance graph, Shape-CFD detects that Article 1092’s sentence-level embedding cloud has high affinity with both the query and a procedural provision about evidence. The convection-diffusion dynamics propagate the relevance signal from the highly-relevant Article 1092 to semantically adjacent provisions, while the advection term ensures directionality: information flows preferentially from more query-aligned to less query-aligned nodes. Result: Article 1092 is promoted to position 1, with the complementary procedural provision correctly at position 2.

Why pure cosine fails: The centroid embedding of Article 1092 suffers from “semantic averaging”—the marriage-law and property-law sentence embeddings, when averaged, produce a centroid that is approximately equidistant from the query as several less relevant provisions. The Chamfer distance, by matching at the sentence level, correctly identifies the strong local alignment.

---

5 Analysis: Physical Interpretation of the Optimal Regime

5.1 Why Weak Advection Outperforms Strong Advection

The empirical finding that $\text{Pe} \approx 0.065$ is optimal—a regime where diffusion dominates advection by a factor of $\sim 15\times$—may appear counterintuitive. If the advection term encodes query direction, why does increasing it degrade performance?

We offer a physical interpretation grounded in transient-state dynamics. The convection-diffusion equation admits two characteristic timescales:

- **Diffusive timescale:** $\tau_D \sim L^2 / D$, where L is the graph diameter
- **Advective timescale:** $\tau_U \sim L / U$, where $U = \overline{|u_{ij}|}$

The Péclet number is their ratio: $\text{Pe} = \tau_D / \tau_U = UL/D$. When $\text{Pe} \ll 1$, diffusion equilibrates much faster than advection can transport concentration directionally. Our early-stopping criterion

(ϵ=10−3\epsilon = 10^{-3}, typically 30–50 iterations) truncates the dynamics well before the advective transport has fully developed. In this transient regime:

1. **Diffusion performs global smoothing:** The fast diffusive mode rapidly propagates the initial cosine-similarity signal across local neighborhoods, correcting isolated misrankings.
2. **Advection provides a directional perturbation:** The slow advective mode contributes a subtle, systematic bias toward query-aligned nodes, improving ranking without overwhelming the stable diffusive backbone.

When Pe\text{Pe} is increased beyond ∼0.1\sim 0.1, the advective timescale becomes comparable to the diffusive timescale, and the early-stopped transient captures unstable advective oscillations before they are damped by diffusion. This manifests as convergence failure (the iteration diverges or oscillates) and degraded ranking quality.

Analogy to numerical methods. This phenomenon parallels the classical observation in CFD that explicit upwind schemes are unconditionally stable for diffusion-dominated flows (Pe⩽1\text{Pe} \leq 1 locally) but require increasingly aggressive CFL restrictions or implicit methods for advection-dominated regimes. Our explicit scheme naturally operates best in the diffusion-dominated regime.

5.2 The Advection Signal in High-Dimensional Embedding Spaces

The curse of dimensionality provides a complementary explanation for weak advection. For vectors uniformly distributed on the dd-dimensional unit sphere, inner products concentrate around zero: ⟨x,y⟩∼𝒩(0,1d)\langle x, y \rangle \sim \mathcal{N}\left(0, \frac{1}{d}\right)

With d=4096d = 4096, the standard deviation is σ≈0.016\sigma \approx 0.016, meaning typical advection coefficients uij∈[−0.048,+0.048]u_{ij} \in [-0.048, +0.048] (within 3σ3\sigma). These values are inherently small compared to the diffusion coefficient D=0.15D = 0.15, making the weak-advection regime the natural operating point of the algorithm in high-dimensional spaces, rather than an arbitrary choice.

This offers a theoretical justification for the algorithm’s robustness: the high dimensionality of modern embedding spaces naturally constrains the system to the stable, diffusion-dominated regime, making the algorithm self-regularizing without requiring careful tuning of the advection strength.

5.3 Limitations

We acknowledge several limitations of the current work, including three fundamental theoretical issues that we identify as priorities for future resolution:

1. **Metric space inconsistency in advection (theoretical).** As discussed in Section 3.6, Shape-CFD’s graph topology is defined by the Chamfer distance (a set-theoretic metric on point clouds), while the advection field is computed from centroid vector differences (a Euclidean metric on single points). This places the diffusion and advection operators in incompatible metric spaces—the graph “knows” about fine-grained sentence matches, but the advection “sees” only the collapsed centroid. This hybrid formulation, while empirically effective, lacks rigorous differential-geometric justification and may fail on documents with highly multimodal sentence distributions.

2. **Noise dilution in symmetric Chamfer distance (practical).** The symmetric Chamfer formulation systematically penalizes long, heterogeneous documents. When a document contains KK sentences but only one is relevant, the (K−1)(K-1) irrelevant sentences contribute large reverse-direction distances that dilute the signal. This contradicts the established asymmetric matching principle validated by ColBERT’s MaxSim operator [22], and may explain Shape-CFD’s limited gains on corpora with predominantly short documents (e.g., FiQA).

3. **JL projection introduces noise, not signal (theoretical).** The O(1/m)O(1/m) variance expansion from Johnson-Lindenstrauss projection is stochastic noise injected by the random matrix, not a meaningful physical signal. The algorithm’s reliance on diffusion dominance (Pe≈0.065\text{Pe} \approx 0.065) effectively marginalizes the advection term—including its noise components—but a PDE framework whose directional transport mechanism is driven by projection artifacts is theoretically fragile. Scalar potential convection (Section 6) would eliminate this dependence entirely.

4. **Scale.** Our experiments use N=30N = 30 candidates. Scaling to N=100N = 100+ would require investigating graph connectivity and the O(N2)O(N^2) similarity matrix cost.

5. **Ground-truth evaluation.** The LawVein results rely on LLM blind evaluation rather than human-annotated relevance judgments.

6. **Over-smoothing without reaction terms.** The linear diffusion dynamics guarantee convergence to a flat equilibrium as t→∞t \to \infty. Our early-

stopping strategy mitigates this practically but does not resolve it theoretically. **v3 update:** Section 4.3.3 demonstrates that the Allen-Cahn reaction term, while theoretically sound, is numerically dormant in the short-query regime ($\Delta C_0 \approx 0.1$) and cannot substitute for early stopping in this operating regime.

7. **Domain-dependent initial field (v3).** Section 4.3.4 demonstrates that the mixed initial field $C_0 = \alpha \cdot \text{MaxSim} + (1-\alpha) \cdot \text{MeanSim}$ yields a +3.2% improvement on NFCorpus at $\alpha = 0.7$ but degrades on SciFact at all $\alpha > 0.5$. This makes the mixed initial field a domain-specific enhancement rather than a universal improvement, limiting its applicability in zero-shot settings where corpus characteristics are unknown.

8. **PDE hurts at token level (v5).** When PDE reranking is applied sequentially on top of token Chamfer scores (using sentence-level graph structure), performance degrades from 0.3214 to 0.3180. This is because the sentence-level graph topology is defined by coarse centroid similarities, while the token-level initial concentrations encode fine-grained per-token matches—the two operate in mismatched semantic spaces. The "sampled graph" variant (constructing a graph from token subsets) produces identical degradation (0.3180), confirming that the mismatch is structural, not merely a resolution artifact. The correct integration strategy is score-level fusion (Section 3.9.2), which treats PDE as an orthogonal signal source rather than a sequential smoother.

9. **Allen-Cahn reaction term is order-preserving (v5 theoretical note).** By the ODE comparison theorem, the Allen-Cahn reaction term $R(C) = \gamma C(1-C)(C-\theta)$ preserves the ordering of concentrations: if $C_i > C_j$ at time t, then $C_i > C_j$ for all future t under the reaction dynamics alone (without diffusion). Since NDCG is a rank-based metric, a monotone transformation of scores cannot change the ranking. This provides a theoretical explanation—independent of the numerical dormancy argument (Section 4.3.3)—for why the reaction term cannot improve ranking quality in isolation.

10. **Embedding model dependency (v5).** All V11 token-level results use Qwen3-Embedding-8B, a model primarily trained on Chinese and English text. The NFCorpus benchmark consists of English biomedical abstracts, which is not the model's primary training domain. The token-level point cloud quality is fundamentally dependent on the embedding model's per-token representation quality; models with stronger per-token

representations (e.g., domain-specific biomedical encoders) may yield different—potentially higher—results. The generalizability of the token-level approach across embedding architectures remains to be validated.

4.3.5 Subspace Chamfer Distance (v3 supplement)

A key question arising from the JL projection analysis (Section 5.3, Limitation 3) is whether alternative dimensionality strategies can improve Chamfer distance discrimination without the noise injection inherent to random projection. We investigate a structured decomposition approach inspired by Product Quantization [Jégou et al., 2011] and the Sparse Grid method from high-dimensional CFD [Bungartz and Griebel, 2004].

Method. We partition the d=4096d = 4096 dimensional embedding space into M=32M = 32 contiguous subspaces of dimension ds=128d_s = 128. For each subspace mm, we compute an independent cosine distance dm(a,b)=1−cos(a[m],b[m])d_m(a, b) = 1 - \cos(a_{[m]}, b_{[m]}), where a[m]a_{[m]} denotes the mm-th 128-dimensional slice of vector aa. The aggregate PQ-Chamfer distance is defined as the mean over subspaces:

dPQ(a,b)=1M∑m=1Mdm(a[m],b[m])d_{\text{PQ}}(a, b) = \frac{1}{M} \sum_{m=1}^{M} d_m(a_{[m]}, b_{[m]})

This decomposition preserves the original dimensional structure—unlike JL projection, which stochastically mixes all dimensions via a random Gaussian matrix. Each 128-dimensional subspace exhibits higher distance variance than the full 4096-dimensional space (the concentration of measure effect weakens as dd decreases), thereby improving the discriminative power of Chamfer distance.

Results (NFCorpus, 323 queries):

Method	NDCG@10	Δ vs V4
Shape-CFD V4 (4096d cosine distance)	0.2508	—
Shape-CFD + JL128 (random projection to 128d)	0.2316	−7.6%
Shape-CFD + PQ-Chamfer (32×128 subspaces)	0.2684	+7.0%
Shape-CFD + PQ-Chamfer (64×64 subspaces)	**0.2735**	**+9.1%**

An ablation study over the subspace granularity reveals a monotonic improvement with finer partitions: 16×256 (0.2650) → 32×128 (0.2684) → 64×64 (0.2735), consistent with the O(1/d)O(1/\sqrt{d}) distance variance scaling. Variance-weighted aggregation (weighting subspaces by query energy) degrades performance by −3.7%, indicating that Qwen3-Embedding distributes semantic information uniformly across dimensions. Dimension permutation has negligible effect (−0.1%), confirming that the embedding dimensions lack local semantic structure.

The PQ-Chamfer decomposition achieves a +9.1% improvement over the full-dimensional baseline, compared to the −7.6% degradation of JL projection—a total swing of 16.7 percentage points. This confirms that the critical factor is not dimensionality reduction per se, but whether the dimensional transformation preserves or destroys the embedding structure.

Physical interpretation. In the CFD analogy, this corresponds to replacing a uniform coarse mesh (full 4096d cosine, low resolution due to concentration) with a multi-block structured mesh (32 independent 128d domains, each with higher effective resolution). The subspace distances act as independent "sensors" that detect fine-grained semantic differences invisible to the full-dimensional metric, analogous to Adaptive Mesh Refinement (AMR) applied to different regions of the embedding space.

5.4 Subspace Virtual Tokens: PQ Decomposition as Semantic Factorization (v4 supplement)

The PQ-Chamfer distance introduced in Section 4.3.5 decomposes the 4096-dimensional embedding space into 64 contiguous 64-dimensional subspaces and computes independent cosine distances within each, yielding a +8.8% improvement over full-dimensional Chamfer. In this v4 supplement, we discover that the **order of operations**—subspace aggregation versus sentence-level Chamfer—critically affects retrieval quality, and that reversing the standard order unlocks an additional +2.5% improvement at zero computational cost.

The Operation Order Insight

Standard PQ-Chamfer (Section 4.3.5) first aggregates across 64 subspaces to produce a single scalar distance between two vectors, then applies the Chamfer operator across sentences: dPQ-Chamfer(q,D)=Chamfer({q},D;dPQ)d_{\text{PQ-Chamfer}}(q, D) = \text{Chamfer}\left(\{q\}, D; \; d_{\text{PQ}}\right)

VT-Aligned reverses this order: for each subspace ss independently, it first computes the sentence-level Chamfer distance (in 64 dimensions), then averages

across subspaces:
$$d_{\text{VT}}(q, D) = \frac{1}{S} \sum_{s=0}^{S-1} \left[\min_{j \in D} d_s(q, s_j) + \frac{1}{|D|} \sum_{j \in D} d_s(s_j, q) \right]$$

where $d_s(a,b) = 1 - \cos(a^{(s)}, b^{(s)})$ is the cosine distance restricted to the s-th 64-dimensional subspace, and $a^{(s)}$ denotes the s-th subspace slice of vector a.

Physical interpretation. Each 64-dimensional subspace acts as an independent "virtual token"—a semantic dimension that captures one facet of the full embedding. The query's single 4096d vector is effectively decomposed into 64 independent semantic probes, each searching for its best match across all document sentences in its own 64-dimensional subspace. This increases the effective query density from 1 point to 64 points without any model re-inference.

We also evaluate an **unaligned** variant (VT-Unaligned) that flattens all subspaces into a single pool of 64d virtual tokens and computes full Chamfer in 64d space, allowing cross-subspace matching (subspace i of the query can match subspace $j \neq i$ of a document sentence).

Stacking with Adjoint Prefetch

We combine VT-Aligned with the V7.1 adjoint state prefetch (Section 5.5), using VT-Aligned distance for graph construction and PDE initial conditions, and the adjoint flux formula for boundary node selection:
$$i^* = \arg\max_{i \in \partial V} C_i \cdot \max\left(0, \sum_{j \in \mathcal{N}(i)} U_{ij} W_{ij} C_j\right)$$

Probe vectors are constructed via cosine interpolation ($\text{probe} = 0.7 \hat{q} + 0.3 \hat{d}_i$) and retrieved via cosine similarity over the full corpus.

Results (NFCorpus, 323 queries)

Method	NDCG@10	vs Cosine	vs PQ-Chamfer
Cosine baseline	0.2195	—	—
PQ-Chamfer (Section 4.3.5)	0.2729	+24.3%	baseline
VT-Unaligned	0.2753	+25.4%	+0.9%
V7.1 Adjoint Prefetch	0.2759	+25.7%	+1.1%
VT-Aligned	**0.2797**	**+27.4%**	**+2.5%**

Method	NDCG@10	vs Cosine	vs PQ-Chamfer
VT-Aligned + V7.1 Adjoint (VT-V7, JS)	0.2844	+29.6%	+4.2%
Rust Pipeline (top_n=55, calibrated)	**0.2852**	**+29.9%**	**+4.5%**

Key findings:

1. **VT-Aligned** > **VT-Unaligned** (+0.9% vs +2.5%): Cross-subspace matching introduces noise, consistent with the LID analysis (Section 4.3.5) showing Qwen3-8B embeddings distribute information uniformly across dimensions.

2. **VT-Aligned** + **V7.1 stack near-orthogonally**: Individual gains are +2.5% and +1.1%; combined gain is +4.2%, exhibiting slight superlinear stacking. This confirms the two innovations address independent bottlenecks (distance granularity vs. candidate pool expansion).

3. **Zero additional cost**: VT-Aligned requires no model re-inference, no additional data, and no new hyperparameters. The only change is the order of mathematical operations on existing sentence vectors.

Rust Acceleration

We implement the full VT-V7 pipeline (centroid cosine pre-filter top_n=55 → VT-Aligned graph construction → PDE solving) as a native Rust module (`law-vexus`), achieving **38 ms per query** on commodity CPU hardware (128 workers, 323 queries in ~27s)—a 130x speedup over the JavaScript reference implementation. The Rust module includes 45 unit tests and uses `rayon` for parallel distance matrix computation on multi-core systems.

Calibration. During Rust re-implementation, we identified and corrected 6 discrepancies with the JavaScript reference: 1. `cosine_distance_64d`: removed `.max(0.0)` clamping that suppressed negative cosine distances 2. `build_knn` edge weight: changed from `1/(1+d)` to `exp(-2d)` to match JS `beta=2.0` 3. `build_knn` graph symmetrization: enforced Wij=WjiW_{ij}=W_{ji} for conservation 4. `compute_advection`: corrected to centroid vector difference projected onto query direction 5. `solve_pde` time step: CFL-adaptive `min(0.1, 0.8/maxDeg)` 6. `cosine_top_n`: changed from sentence-level MaxSim to document centroid cosine for pre-filtering

After these corrections, the Rust engine achieves NDCG@10 = **0.2852** (+29.9% vs cosine), surpassing the JavaScript reference (0.2844). The calibration trajectory:

0.2145 (initial) → 0.2597 (formula alignment) → 0.2750 (centroid pre-filter) → 0.2852 (top_n=55 optimal).

5.5 The Pareto Boundary of Initial Fields (v3 supplement)

The dataset-dependent behavior of the mixed initial field (Section 4.3.4) admits a precise statistical interpretation. We formalize this as the **Dual-Role Theorem of MeanSim**.

Theorem 2 (Dual-Role Theorem). In the mixed initial field C0=α·MaxSim+(1−α)·MeanSimC_0 = \alpha \cdot \text{MaxSim} + (1-\alpha) \cdot \text{MeanSim}, the MeanSim component plays two distinct roles depending on document structure:

(i) Low-pass filter (short, homogeneous documents). When Ki≤KcritK_i \leq K_{\text{crit}} (few sentences, all topically coherent), MaxSim≈MeanSim\text{MaxSim} \approx \text{MeanSim} and the decoupling adds no semantic information. However, MaxSim introduces extreme-value sampling noise (variance ∼1/Ki\sim 1/K_i), while MeanSim provides variance reduction via averaging. Increasing α\alpha degrades performance monotonically by amplifying noise without signal gain.

(ii) Macroscopic gravity (long, heterogeneous documents). When Ki≫KcritK_i \gg K_{\text{crit}} (many sentences with sparse relevance), MeanSim suffers catastrophic signal dilution (MeanSim→0\text{MeanSim} \to 0 as Ki→∞K_i \to \infty for fixed relevant sentences). MaxSim penetrates this dilution by isolating the highest-scoring sentence. The optimal α*\alpha^* balances the signal gain from MaxSim against its false-positive risk, yielding the inverted-U response observed on NFCorpus.

The transition between regimes is governed by the effective signal-to-noise ratio of the MaxSim operator:

SNRMaxSim=𝔼[maxscos(q,srel)]−𝔼[maxscos(q,sirrel)]Var[maxscos(q,s)]1/2\text{SNR}_{\text{MaxSim}} = \frac{\mathbb{E}[\max_s \cos(q, s_{\text{rel}})] - \mathbb{E}[\max_s \cos(q, s_{\text{irrel}})]}{\text{Var}[\max_s \cos(q, s)]^{1/2}}

When SNRMaxSim>1\text{SNR}_{\text{MaxSim}} > 1 (long documents with strong local evidence), decoupling is beneficial; when SNRMaxSim<1\text{SNR}_{\text{MaxSim}} < 1 (short documents where the extreme value is noise-dominated), the symmetric Chamfer baseline (α=0.5\alpha = 0.5) is optimal.

Scale-adaptive symmetry breaking. Based on this analysis, we propose a document-level adaptive mixing coefficient:

αi=0.5+Δmax·tanh(max(0,|Di|−LminLscale))\alpha_i = 0.5 + \Delta_{\max} \cdot \tanh\left(\max\left(0, \frac{|D_i| - L_{\min}}{L_{\text{scale}}}\right)\right)

where |Di||D_i| is the sentence count of document ii, LminL_{\min} is the minimum length below which α=0.5\alpha = 0.5 (the Chamfer default), LscaleL_{\text{scale}} controls the transition steepness, and Δmax∈[0.2,0.25] \Delta_{\max} \in [0.2, 0.25] caps the maximum asymmetry. This formula satisfies two critical properties: (1) it reduces to the v1 Chamfer baseline (α≡0.5\alpha \equiv 0.5) for corpora of uniformly short documents, ensuring no degradation; (2) it smoothly transitions to the α=0.7\alpha = 0.7–0.750.75 regime for long-document corpora, capturing the NFCorpus-type improvement. Experimental validation of this adaptive formula is deferred to future work.

5.5 Speculative Retrieval: From Reranker to Retriever (preliminary)

A fundamental limitation of reranking-based approaches is the **candidate ceiling**: the reranker can only reorder documents already present in the initial retrieval set. If a relevant document is ranked beyond the top-NN cutoff by the initial cosine retriever, no amount of reranking can recover it.

We propose **Speculative Retrieval**, inspired by CPU branch prediction: after an initial PDE pass on the top-NN candidates, we identify high-pressure boundary nodes—documents with high post-PDE concentration but low graph degree (i.e., at the "edge" of the semantic neighborhood). These nodes indicate that the candidate set boundary intersects a region of high relevance, and that additional relevant documents may exist just beyond the retrieval horizon.

For each high-pressure boundary node, we compute a diffusion probe vector by interpolating between the node's embedding and the query direction, then retrieve the nearest neighbors of this probe from the full corpus. The expanded candidate set is then subjected to a second PDE pass.

Preliminary results (NFCorpus, 323 queries):

Method	**NDCG@10**	**Δ vs PQ64-top30**
PQ-Chamfer top-30 (baseline)	0.2729	—
Naive expansion to top-40	0.2766	+1.4%
Naive expansion to top-50	0.2757	+1.0%
Speculative Retrieval	**0.2781**	**+1.9%**

The speculative approach outperforms naive pool expansion (+1.9% vs +1.4%), confirming that the PDE concentration field provides directional information

beyond what cosine similarity alone captures. This represents a preliminary step toward upgrading Shape-CFD from a pure reranker to a retrieval-augmented reranker that can actively discover relevant documents missed by the initial retriever.

6 Conclusion

We have presented Shape-CFD, a physics-inspired reranking framework that models post-retrieval reranking as a conservative convection-diffusion process on a document similarity graph. By introducing directional advection alongside isotropic diffusion, and by upgrading document representations from single-point embeddings to sentence-level point clouds measured via Chamfer distance, Shape-CFD captures both inter-document structural relationships and fine-grained sub-document semantic matches.

v5 supplement findings. The V11 token-level point cloud extension (Section 3.9) represents a paradigm-level advance: by extracting per-token last hidden states rather than pooled sentence embeddings, each document becomes a dense point cloud of $$356 tokens and each query unfolds into $$6 token vectors. The Token PQ-Chamfer distance—a zero-parameter hard cross-attention mechanism—achieves NDCG@10 = 0.3214 (+46.4% vs cosine) on NFCorpus via full scan. A two-stage centroid acceleration pipeline (Section 3.9.1) achieves 0.3220 (+46.7%) in only 23 ms with lossless accuracy, demonstrating that token centroid pre-filtering is both efficient and noise-reducing. Most significantly, the discovery that PDE provides an orthogonal signal to token Chamfer (Section 3.9.2)—local nearest-neighbor matching vs. global semantic propagation—enables score-level fusion that achieves the overall best NDCG@10 = **0.3232** (+47.2% vs cosine) at 26 ms latency on CPU, with zero cross-encoder inference and zero training. The complete 10-method ablation validates each component’s contribution and establishes clear diminishing-returns boundaries.

v4 supplement findings. The VT-Aligned virtual token decomposition (Section 5.4) demonstrates that PQ subspace structure can be repurposed as a semantic factorization tool: by computing sentence-level Chamfer within each 64-dimensional subspace independently, the effective query representation density increases 64-fold without model re-inference. Combined with the adjoint state prefetch, the VT-V7 configuration achieves NDCG@10 = 0.2844 (+29.6% over cosine) in JavaScript. A calibrated Rust full-pipeline engine with 6 bug fixes and centroid-based cosine pre-filter (top_n=55) further improves to NDCG@10 = **0.2852** (+29.9% over cosine). The Rust implementation with SIMD-friendly memory layout and `rayon` parallelism reduces end-to-end latency to 38 ms per query—a 130x speedup over the JavaScript reference.

v3 supplement findings. The systematic ablation of V5 reaction-diffusion extensions (Section 4.3.3) reveals that the Allen-Cahn reaction term is physically dormant in the short-query, narrow score-range regime—not due to implementation deficiency, but as a consequence of the cubic scaling of reaction forces versus linear scaling of diffusion in the $\Delta C_0 \ll 1$ limit. The mixed initial field experiment (Section 4.3.4) demonstrates that decoupling the symmetric Chamfer distance into explicit MaxSim and MeanSim components yields a domain-dependent response: a +3.2%+3.2\% improvement on NFCorpus (heterogeneous long documents) at $\alpha=0.7$, but monotonic degradation on SciFact (homogeneous short documents). This Pareto boundary (Section 5.4) is governed by the trade-off between extreme-value signal gain and sampling noise, and the original symmetric Chamfer ($\alpha=0.5$) remains the Pareto-optimal default for domain-agnostic deployment.

Future Work

The three theoretical limitations identified in Section 5.3 point directly to a coherent next-generation design. We outline the key directions:

1. **Scalar potential convection (resolves Limitations 1 and 3).** The metric space inconsistency and JL noise dependence share a common root cause: defining advection via vector inner products in high-dimensional space. We propose replacing the centroid-based advection with a scalar potential gradient, following the Darcy flow analogy: $u_{ij}^{\text{potential}} = \kappa \cdot (\Phi_j - \Phi_i)$ where $\Phi_i = \max_{s \in \mathcal{P}_i} \cos(q, s)$ is the MaxSim score of document i against the query. This formulation eliminates the curse of dimensionality entirely—scalar subtraction has no dependence on the ambient dimension d—and unifies the metric space: both the graph topology and the advection field operate on the same score manifold. The resulting advection is guaranteed to be non-degenerate whenever two documents have different MaxSim scores, regardless of the embedding dimension.

2. **Allen-Cahn reaction-diffusion for long-document retrieval (resolves Limitation 6, v3 update).** Section 4.3.3 empirically confirms that the reaction term is dormant in the short-query regime ($\Delta C_0 \approx 0.1$). However, in long-document retrieval scenarios (e.g., legal case matching, multi-hop QA), where $\Delta C_0 \gg 0.1$ due to sparse evidence distributed across lengthy documents, the Allen-Cahn dynamics are expected to activate meaningfully. The domain-adapted formulation

γeff=γbase/(ΔC)2\gamma_{\text{eff}} = \gamma_{\text{base}} / (\Delta C)^2 (Section 4.3.3) provides the correct scaling law for such scenarios. We propose augmenting the convection-diffusion equation with the domain-adapted reaction term: R(Ci)=γeff(Ci−Cmin)(Cmax−Ci)(Ci−θOtsu)R(C_i) = \gamma_{\text{eff}} \, (C_i - C_{\min})(C_{\max} - C_i)(C_i - \theta_{\text{Otsu}}) where θOtsu\theta_{\text{Otsu}} is dynamically determined via Otsu’s method. Target benchmarks include COLIEE (legal case matching) and TREC-Robust04 (narrative long queries).

3. **Asymmetric shape matching (resolves Limitation 2).** Replacing the symmetric Chamfer distance with a directed, query-weighted variant: dasym(𝒫q,𝒫i)=1|𝒫q|∑a∈𝒫qminb∈𝒫idcos(a,b)d_{\text{asym}}(\mathcal{P}_q, \mathcal{P}_i) = \frac{1}{|\mathcal{P}_q|} \sum_{a \in \mathcal{P}_q} \min_{b \in \mathcal{P}_i} d_{\cos}(a, b) This measures only whether the query’s semantic facets are covered by the document, without penalizing unmatched document sentences. Combined with LSH (locality-sensitive hashing) sign quantization and hardware popcount instructions, this could achieve sub-millisecond shape distance computation.

4. **Spectral acceleration.** For larger candidate sets (N≫100N \geq 100), evolving the concentration field in the spectral domain of the graph Laplacian could reduce the iteration count from O(N)O(N) to O(logN)O(\log N).

5. **Scale-adaptive initial field (resolves Limitation 7, v3).** The Pareto boundary analysis (Section 5.4) motivates a document-level adaptive mixing coefficient αi=0.5+Δmax⋅tanh(max(0,(|Di|−Lmin)/Lscale))\alpha_i = 0.5 + \Delta_{\max} \cdot \tanh(\max(0, (|D_i| - L_{\min}) / L_{\text{scale}})) that automatically transitions from the symmetric Chamfer baseline (α=0.5\alpha = 0.5) for short documents to the asymmetric MaxSim-dominant regime (α≈0.7\alpha \approx 0.7) for long documents. Combined with per-query calibration via the MaxSim SNR criterion, this would yield a fully adaptive, domain-agnostic initial field.

6. **Chunk-level graph construction for long-document retrieval.** For corpora with documents exceeding hundreds of sentences (e.g., legal case files, clinical trial reports), document-level graph construction loses the fine-grained Riemannian manifold structure that underpins Shape-CFD’s success. We propose segmenting long documents into semantic chunks and constructing the KNN graph at the chunk level, creating a much larger but topologically richer graph where the Allen-Cahn reaction term (item 2) is expected to become physically active due to the expanded score range (ΔC0≫0.1\Delta C_0 \gg 0.1).

7. **Token-level graph construction (resolves v5 Limitation 8).** The current fusion approach treats PDE and token Chamfer as independent signals combined at the score level. A more principled integration would construct the KNN graph directly from token-level distances (e.g., using Token PQ-Chamfer as the edge weight metric), enabling PDE to operate in the same semantic space as the token-level initial concentrations. The computational challenge is the O(N2·M2)O(N^2 \cdot M^2) cost of pairwise token Chamfer between all candidate documents; approximation via token centroid graphs or locality-sensitive hashing may make this tractable.

8. **Cross-embedding generalization (resolves v5 Limitation 10).** The token-level pipeline’s dependence on Qwen3-Embedding-8B should be validated across diverse embedding architectures (e.g., BGE-M3, E5-Mistral, domain-specific biomedical encoders). The key question is whether the per-token hidden state quality varies significantly across architectures, and whether the optimal PQ subspace granularity (S=64S = 64) is architecture-dependent.

9. **Learned fusion weights.** The current fusion weight λ=0.7\lambda = 0.7 is determined by grid search on NFCorpus. A lightweight learned fusion (e.g., per-query λ\lambda predicted from query token statistics such as token count, centroid norm, or inter-token variance) could adapt the fusion balance to query characteristics without requiring cross-encoder-scale computation.

References

[1] L. Vilnis and A. McCallum. Word representations via Gaussian embedding. In Proc. ICLR, 2015.

[2] S. Dampanaboina, J. Shen, and B. Khoshnevisan. Diffusion-Aided Joint Source Channel Coding for Retrieval-Augmented Generation. arXiv preprint, 2024.

[3] N. Lin, Z. Bai, Y. Chen, and S. Lin. Advection augmented convolutional neural networks. In Proc. IEEE/CVF CVPR, 2024.

[4] S. Brin and L. Page. The anatomy of a large-scale hypertextual web search engine. Computer Networks, 30(1-7):107–117, 1998.

[5] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, and A. Anandkumar. Fourier neural operator for parametric partial differential equations. In Proc. ICLR, 2021.

[6] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. Journal of Computational Physics, 378:686–707, 2019.

[7] W. Johnson and J. Lindenstrauss. Extensions of Lipschitz mappings into a Hilbert space. Contemporary Mathematics, 26:189–206, 1984.

[8] J. Tremblay, S. Prakash, D. Acuna, M. Brophy, V. Jampani, C. Anil, T. To, E. Cameracci, S. Bonn, and S. Birchfield. Training deep networks with synthetic data: Bridging the reality gap by domain randomization. In Proc. CVPR Workshops, 2018.

[9] R.I. Kondor and J. Lafferty. Diffusion kernels on graphs and other discrete structures. In Proc. ICML, 2002.

[10] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. Yih, T. Rocktäschel, S. Riedel, and D. Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In Proc. NeurIPS, 2020.

[11] L. Gao, X. Ma, J. Lin, and J. Callan. Tevatron: An efficient and flexible toolkit for neural retrieval. arXiv preprint arXiv:2203.05765, 2022.

[12] F. Chung. The heat kernel as the pagerank of a graph. Proc. National Academy of Sciences, 104(50):19735–19740, 2007.

[13] B. Chamberlain, J. Rowbottom, M. Gorinova, M. Bronstein, S. Webb, and E. Rossi. GRAND: Graph neural diffusion. In Proc. ICML, 2021.

[14] K. Xu, W. Hu, J. Leskovec, and S. Jegelka. How powerful are graph neural networks? In Proc. ICLR, 2019.

[15] G. Fantuzzi, D. Goluskin, D. Huang, and S.I. Chernyshenko. Bounds for deterministic and stochastic dynamical systems using sum-of-squares optimization. SIAM Journal on Applied Dynamical Systems, 15(4):1962–1988, 2016.

[16] M. Defferrard, X. Bresson, and P. Vandergheynst. Convolutional neural networks on graphs with fast localized spectral filtering. In Proc. NeurIPS, 2016.

[17] I. Oseledets. Tensor-train decomposition. SIAM Journal on Scientific Computing, 33(5):2295–2317, 2011.

[18] S. Kim, J. Lee, and M. Seo. G-RAG: Knowledge-graph augmented language models for improved retrieval-augmented generation. arXiv preprint, 2024.

[19] L. Chen and Y. Liu. Graph-augmented dense statute retrieval for legal judgment prediction. In Proc. ACL, 2023.

[20] R. Huang, N. Marin, and J. Yang. Spreading activation for knowledge graph-augmented retrieval. In Proc. EMNLP, 2024.

[21] D. Mussmann, S. Jain, and S. Ermon. Gaussian embeddings for large-scale document retrieval. In Proc. NeurIPS, 2024.

[22] O. Khattab and M. Zaharia. ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In Proc. SIGIR, 2020.

[23] B. Eliasof, E. Haber, and E. Treister. PDE-GCN: Novel architectures for graph neural networks motivated by partial differential equations. In Proc. NeurIPS, 2021.

[24] T.N. Kipf and M. Welling. Semi-supervised classification with graph convolutional networks. In Proc. ICLR, 2017.

[25] N. Thakur, N. Reimers, A. Rücklé, A. Srivastava, and I. Gurevych. BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. In Proc. NeurIPS Datasets and Benchmarks, 2021.